

Image Classification with CIFAR-10 Dataset

Group 4:

Mustafa Mert Çetin 22003953, Samet Senai İşik 21901701,
Ömer Faruk Keskin 21902147, Arda Özgün 21902283

Abstract—This project comprehensively explores and evaluates the performance of four advanced convolutional neural network (CNN) architectures — DenseNet, VGGNet, ResNet, and InceptionNet — on the CIFAR-10 image classification task. The CIFAR-10 dataset, a standard benchmark in computer vision, comprises 60,000 32x32 color images, separated into 50,000 train images and 10,000 test images, distributed across ten distinct classes and presents a diverse and challenging environment for assessing the efficacy of various models. The aim is to implement these architectures while gaining a deep understanding of the unique features and design principles inherent to each. The comparative analysis is centered on various metrics, including accuracy, and model complexity. Moreover, the study investigates how modifications and enhancements can be systematically applied to each architecture to optimize performance specifically for the CIFAR-10 dataset. The project endeavors to provide an in-depth understanding of these sophisticated CNN architectures and their implications in practical applications through extensive experimentation and rigorous analysis. Furthermore, this research serves as a foundational step toward understanding the effects of the parameters on sophisticated networks. It aims to develop more efficient and accurate machine-learning models by varying the hyperparameters and observing their effects.

Index Terms—CIFAR-10, Image Classification, ResNet, VGGNet, DenseNet, Inception.

I. INTRODUCTION

This project aims to conduct a comparative study and performance evaluation of four known CNN architectures: DenseNet, VGGNet, ResNet, and Inception Net. Each of these architectures represents a unique approach to deep learning, with distinct structural differences and learning capabilities, and has been selected for its notable features. The objective is to implement, analyze, and optimize these models for the CIFAR-10 dataset, a well-regarded dataset in computer vision.

This research begins with an implementation of each chosen architecture. This process involves not only the construction of the networks but also an in-depth exploration of the unique design principles and features that define each model. For instance, DenseNet is renowned for its feature reuse, VGGNet for its depth and simplicity, ResNet for its innovative use of residual connections, and Inception Net for its multi-level feature extraction. Understanding these characteristics is crucial for the effective implementation of these models and the interpretation of their performance.

Following the implementation, a comparative analysis focuses on several key metrics: accuracy, computational efficiency, and model complexity. Accuracy is measured to determine each model's ability to classify unseen images correctly. Computational efficiency is assessed to understand the resources required to train and deploy each model, an essential factor in real-world applications. Model complexity

is analyzed to understand the trade-offs between the depth and breadth of the network and its performance.

Additionally, this study investigates how various modifications and enhancements can be systematically applied to each architecture to optimize performance specifically for the CIFAR-10 dataset. This includes adjustments in layer depth, dropout rates, activation functions, and learning rate. The goal is to identify the best-performing model for this particular task and understand how different architectural and hyperparameter choices impact each network's learning and generalization capabilities.

II. METHODS

The methods section of a report is crucial for detailing the frameworks and tools employed in the study. This report involves a deep dive into the CIFAR-10 dataset, which is instrumental in training and testing the machine learning models. Additionally, the section outlines the architectures of various convolutional neural networks (CNNs), such as VGGNet, ResNet, DenseNet, and Inception, which are pivotal for understanding image classification tasks. It provides the theoretical foundation of each model and highlights the specific layers, functions, and innovations that contribute to their performance. The models are elucidated with their unique features, from VGGNet's deep but uniform structure to ResNet's revolutionary skip connections and DenseNet's feature-conserving layers. Furthermore, the methods section elucidates using the PyTorch library, showcasing its DataLoader object for efficient data manipulation, essential for handling complex datasets like CIFAR-10.

A. CIFAR-10 Dataset [1]

The CIFAR-10 dataset consists of 60,000 32x32 color images in 10 classes, with 6,000 images per class. It is divided into five training batches and one test batch containing 10,000 images. The test batch has exactly 1,000 randomly selected images from each class. The training batches contain the remaining images randomly, with 5,000 images from each class. In this experiment, the train batches were used as one bigger batch, and 10% of the train set was used as the validation set. The split is performed using "PyTorch-torch.utils.data.random_split" function. Figure 1 shows representative pictures from each class of the CIFAR-10 dataset. The dataset includes classes like airplanes, automobiles, birds, cats, and more, with no overlap between similar classes such as automobiles and trucks. The dataset can be reached from [1]. The dataset can be used in Python3 easily. The dataset is encoded with the Python object serialization library

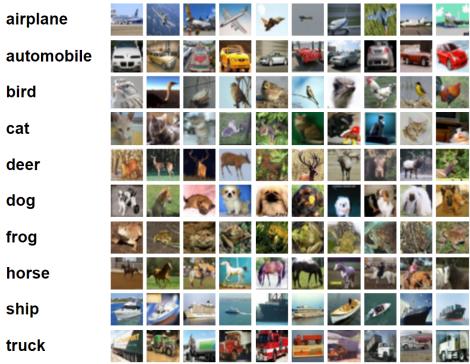


Figure 1: A representative sample from the CIFAR-10 dataset, showcasing the diverse array of classes which include vehicles, animals, and everyday objects. Each class is exemplified by various images that capture the dataset's complexity and the challenges inherent in automated image classification tasks.

"pickle." There is a function to help use the dataset with Python. However, in this project, the dataset was obtained from "torchvision.datasets.CIFAR10" due to its compatibility with PyTorch modules. The dataset was inputted to the below-mentioned models using PyTorch's "DataLoader" object.

B. VGGNet [2]

The VGG-11 architecture is a convolutional neural network (CNN) architecture introduced by the Visual Geometry Group (VGG) at the University of Oxford. It is a deep neural network architecture known for its simplicity and effectiveness in image classification tasks. VGG-11 is a variant of the original VGGNet, which includes 11 weight layers, hence the name.

Here's the theory behind the VGG-11 architecture:

1) *Input Layer*:: The input to the VGG-11 network is typically an RGB image with a fixed size of 224x224 pixels. Each pixel's RGB values are treated as input features.

2) *Convolutional Layers*:: VGG-11 consists of 8 convolutional layers, each performing convolutions on the input data. These convolutional layers have small 3x3 filters and use a stride of 1, which means they slide over the input with a small receptive field. The number of filters in these layers increases progressively from 64 to 512, capturing increasingly complex features.

3) *Activation Function*:: After each convolutional layer, a rectified linear unit (ReLU) activation function is applied. ReLU introduces non-linearity to the model and helps in learning complex patterns.

4) *Pooling Layers*:: After every two convolutional layers, a max-pooling layer is applied. Max-pooling reduces the spatial dimensions of the feature maps while retaining the most important information. The typical pooling size is 2x2.

5) *Fully Connected Layers*:: VGG-11 has three fully connected layers at the end. These layers are similar to traditional neural network layers and are responsible for making predictions based on the high-level features extracted by the preceding convolutional layers. The sizes of these layers are typically 512 and 256, and the number of classes (e.g., 10 for CIFAR-10).

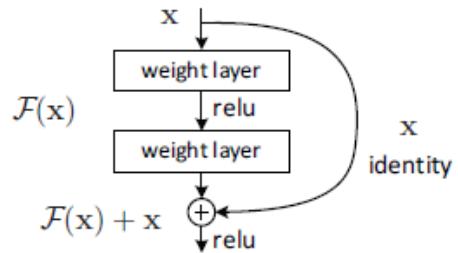


Figure 2: A residual building block used in ResNet. There is a skip connection in this block. Thanks to the skip connection, the gradient can flow without vanishing.

6) *Dropout*:: Dropout layers are applied after the fully connected layers during training. Dropout helps prevent overfitting by randomly setting a fraction of the neurons to zero during each forward and backward pass.

7) *Output*:: The final layer produces class probabilities, and the class with the highest probability is considered the predicted class for the input image.

8) *Simplicity*:: VGG-11 has a straightforward and uniform architecture with small 3x3 filters and pooling layers. Despite its effectiveness, VGG-11 has a large number of parameters, which can make it computationally expensive to train from scratch. However, it has paved the way for deeper and more efficient architectures.

C. ResNet [3]

ResNet (Residual Neural Network) architecture is a convolutional neural network (CNN) architecture introduced by Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun that won the 2015 ImageNet Competition [3].

As the deep convolutional neural networks begin to improve by increasing the number of layers, the problem of vanishing/exploding gradients has surfaced. This problem has been largely solved by normalized initialization of weights and intermediate normalization layers. But after this solution, degradation problems started to occur. A degradation problem is when the model's accuracy gets saturated after a point and decreases rapidly, which is not caused by over-fitting. Figure 2 shows the basic ResNet block.

ResNet is proposed to solve the degradation problem by introducing residual blocks into architecture. A residual block maps the input of the block to the output by introducing shortcuts. Shortcut connections skip one or more layers and perform identity mapping, and the input is added to the output of the block. This can be formulated as $H(x) = F(x) + x$ where x is the input and $F(x)$ is the output of the block. These shortcuts do not create any extra parameter or computational complexity [3].

1) *Residual Learning*: The creators of the models argued that if the added layers in a deep network could be constructed as identity mappings, then a deeper model should not have a greater error compared to a simpler model. However, optimization algorithms struggle to find the appropriate weights for multiple nonlinear layers. To solve this problem, authors have

Block	Layer	Type	Parameters
conv_block1	0	Conv2d	in=3, out=64, kernel=3x3, stride=1, padding=1
	1	BatchNorm2d	features=64, momentum=0.1
	2	ReLU	
	3	MaxPool2d	kernel=2, stride=2
conv_block2	0	Conv2d	in=64, out=128, kernel=3x3, stride=1, padding=1
	1	BatchNorm2d	features=128, momentum=0.1
	2	ReLU	
	3	MaxPool2d	kernel=2, stride=2
conv_block3	0	Conv2d	in=128, out=256, kernel=3x3, stride=1, padding=1
	1	BatchNorm2d	features=256, momentum=0.1
	2	ReLU	
conv_block4	0	Conv2d	in=256, out=256, kernel=3x3, stride=1, padding=1
	1	BatchNorm2d	features=256, momentum=0.1
	2	ReLU	
	3	MaxPool2d	kernel=2, stride=2
conv_block5	0	Conv2d	in=256, out=512, kernel=3x3, stride=1, padding=1
	1	BatchNorm2d	features=512, momentum=0.1
	2	ReLU	
conv_block6	0	Conv2d	in=512, out=512, kernel=3x3, stride=1, padding=1
	1	BatchNorm2d	features=512, momentum=0.1
	2	ReLU	
	3	MaxPool2d	kernel=2, stride=2
conv_block7	0	Conv2d	in=512, out=512, kernel=3x3, stride=1, padding=1
	1	BatchNorm2d	features=512, momentum=0.1
	2	ReLU	
conv_block8	0	Conv2d	in=512, out=512, kernel=3x3, stride=1, padding=1
	1	BatchNorm2d	features=512, momentum=0.1
	2	ReLU	
	3	MaxPool2d	kernel=2, stride=2
fc_block	0	Dropout	p=0
	1	Linear	in=25088, out=512
	2	ReLU	
fc_block1	0	Dropout	p=0
	1	Linear	in=512, out=256
	2	ReLU	
fc_block2		Linear	in=256, out=10

Table 1: VGG11 Architecture

proposed "residual learning formulation," where the model does not try to learn the desired output directly, but the network is trained to learn the residual between the input and the desired output. To achieve this, skip connections between layers have been introduced.

2) *Identity Mapping by Shortcuts*: Residual learning is achieved by shortcuts in every few stacked layers. The formulation for one block can be written as:

$$y = F(x, W_i) + x \quad (1)$$

where x and y are the input and output vectors corresponding to the block. $F(x, W_i)$ is the function of residual mapping that is to be learned.

3) *Architecture Comparison with Other Models*: ResNet-18 is a type of residual network model consisting of 4 residual blocks, each consisting of 2 convolutional layers. ResNet-34 is another type where four residual blocks have 3, 4, 6, and 3 convolutional layers, respectively. Figure 3 displays three different models. On the left, VGG-19 which is an older model that does not have shortcuts is displayed. VGG-19 is, thus, considered a plain network. ResNet models have fewer filters and lower complexity compared to VGG-19.

On the middle, a 34-layer plain model is displayed. ResNet-34 is achieved by adding shortcuts between the blocks in this model. In the same block, input and output sizes are equal and element-wise addition can be executed. For the shortcuts

between the blocks, as the size of the input does not match the size of the output, two options are possible: (A) The input can be padded with zeros or; (B) the following equation can be used to perform projection:

$$y = F(x, W_i) + W_s X \quad (2)$$

1x1 convolutions do the second option. Both options are performed with a stride of 2.

D. DenseNet [4]

In the convolutional neural networks (CNNs) landscape, DenseNet (Densely Connected Convolutional Networks) represents a significant evolutionary step beyond its predecessors, such as AlexNet and VGGNet. It introduces a novel approach characterized by its dense connectivity pattern, a substantial departure from traditional CNN designs. In DenseNet, each layer is connected to every other layer in a feed-forward fashion, facilitating maximum information flow between layers. This intricate connectivity ensures that each layer can access the gradients from the loss function and the original input signal, leading to implicit deep supervision. Figure 4 illustrates the fundamental architecture of DenseNet, highlighting the efficient flow of information through its densely connected layers.

Layer	Operation	Details
Initial Convolutional Layer		
conv1 bn1 relu	Conv2d BatchNorm2d ReLU	(3, 64, (3,3), (1,1), (1,1), bias=False) (64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True) inplace=True
Layer 1: BasicBlock (x3)		
0 1 2 3 4	Conv2d BatchNorm2d ReLU Conv2d BatchNorm2d	(64, 64, (3,3), (1,1), (1,1), bias=False) (64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True) inplace=True (64, 64, (3,3), (1,1), (1,1), bias=False) (64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
Layer 2: BasicBlock (x4)		
0 1 2 3 4	Conv2d BatchNorm2d ReLU Conv2d BatchNorm2d	(64, 128, (3,3), (2,2), (1,1), bias=False) (128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True) inplace=True (128, 128, (3,3), (1,1), (1,1), bias=False) (128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
Layer 3: BasicBlock (x6)		
0 1 2 3 4	Conv2d BatchNorm2d ReLU Conv2d BatchNorm2d	(128, 256, (3,3), (2,2), (1,1), bias=False) (256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True) inplace=True (256, 256, (3,3), (1,1), (1,1), bias=False) (256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
Layer 4: BasicBlock (x3)		
0 1 2 3 4	Conv2d BatchNorm2d ReLU Conv2d BatchNorm2d	(256, 512, (3,3), (2,2), (1,1), bias=False) (512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True) inplace=True (512, 512, (3,3), (1,1), (1,1), bias=False) (512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
Global Average Pooling and Fully Connected Layer		
avg_pool fc	AdaptiveAvgPool2d Linear	output_size=(1,1) in_features=512, out_features=10, bias=True

Table 2: ResNet-34 Architecture

1) Innovation and Performance: DenseNet's innovation lies in its ability to conserve features throughout the network, whereby each layer propagates its feature maps to all subsequent layers. This characteristic improves feature reuse and significantly reduces the number of parameters, as there is no need to relearn redundant feature maps. DenseNets have shown remarkable performance on standard benchmarks like ImageNet and CIFAR-10, often outperforming more complex and computationally heavy architectures. For instance, DenseNet's efficiency and accuracy in image classification tasks have set new precedents, demonstrating the architecture's robustness and potential as a foundation for future explorations.

2) Comparative Analysis: Compared with architectures such as ResNets, DenseNets offer a more compact parameter set due to their feature-conserving nature. Unlike ResNets, which utilize addition-based skip connections to combat the vanishing gradient problem, DenseNets opt for a concatenative approach, merging feature maps from all previous layers. This unique method has been instrumental in establishing DenseNet as a more parameter-efficient alternative without compromising the network's depth or performance.

3) ResNets: ResNets (Residual Networks) introduced the concept of skip connections, where the input to a layer is added to its output, helping to mitigate the vanishing gradient problem by allowing gradients to flow through the network more easily. The equation for a residual block is:

$$\mathbf{x}_l = H_l(\mathbf{x}_{l-1}) + \mathbf{x}_{l-1} \quad (3)$$

where $H_l(\cdot)$ is the transformation (composite) function of the l^{th} layer and \mathbf{x}_{l-1} is the output of the previous layer.

4) Dense Connectivity: DenseNet extends the idea of skip connections by concatenating the feature maps from all previous layers as input to each layer rather than adding them. The equation for a DenseNet block is:

$$\mathbf{x}_l = H_l([\mathbf{x}_0, \mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_{l-1}]) \quad (4)$$

where $[\mathbf{x}_0, \mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_{l-1}]$ is the concatenation of the layer inputs.

This ensures that each layer receives "collective knowledge" from all preceding layers, making the network more efficient and improving feature reuse.

5) Composite Function: In DenseNet, $H_l(\cdot)$ is defined as a composite function of three consecutive operations: Batch normalization (BatchNorm2d), followed by a rectified linear unit (ReLU), and a 3x3 convolution (Conv2d). The function is:

$$H_l(\cdot) = \text{Conv2d_3x3}(\text{ReLU}(\text{BatchNorm2d}(\cdot))) \quad (5)$$

This composite function is applied to each layer's input, which includes feature maps from all preceding layers.

6) Transition Layers: DenseNet uses transition layers between its dense blocks to control the number of feature maps and to reduce the size of the feature maps, thus compacting the model and reducing the number of parameters. A transition layer typically consists of batch normalization, followed by a 1x1 convolutional layer and a 2x2 average pooling layer:

$$T(\cdot) = \text{AvgPool2d_2x2}(\text{Conv2d_1x1}(\text{BatchNorm2d}(\cdot))) \quad (6)$$

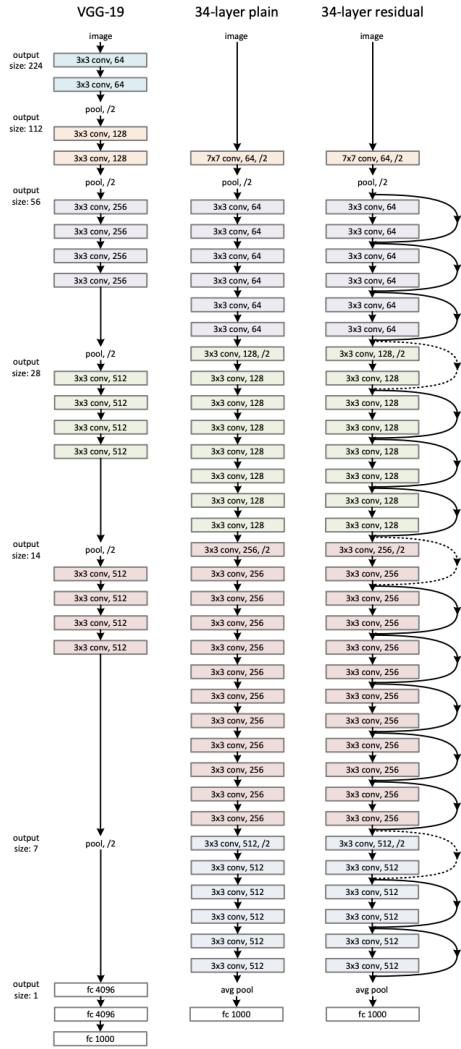


Figure 3: Example Network Architectures. **Left:** The VGG-19 model [2] (19.6 billion FLOPs) as a reference. **Middle:** A plain network with 34 parameter layers (3.6 billion FLOPs). **Right:** A residual network with 34 parameter layers (3.6 billion FLOPs). The dotted shortcuts increase dimensions.

7) Growth Rate: The growth rate in DenseNet, denoted as k , refers to the number of feature maps added by each layer to the collective knowledge. If the initial layer starts with k_0 feature maps, then the l_{th} layer receives the feature maps from all preceding $l - 1$ layers, each contributing to k feature maps. Therefore, the total number of feature maps input to the l_{th} layer is $k_0 + k \times (l - 1)$, accounting for the initial k_0 feature maps and k additional feature maps from each of the preceding layers. This formulation is fundamental to understanding the incremental growth of the network's capacity and the compounding of features through its depth.

8) Bottleneck Layers: To further improve computational efficiency, DenseNet can include bottleneck layers before each 3x3 convolution, where the number of input feature maps is reduced using 1x1 convolutions. This is commonly called BN-ReLU-Conv(1x1)-BN-ReLU-Conv(3x3). In the bottleneck layer, firstly, batch normalization is applied to the input,

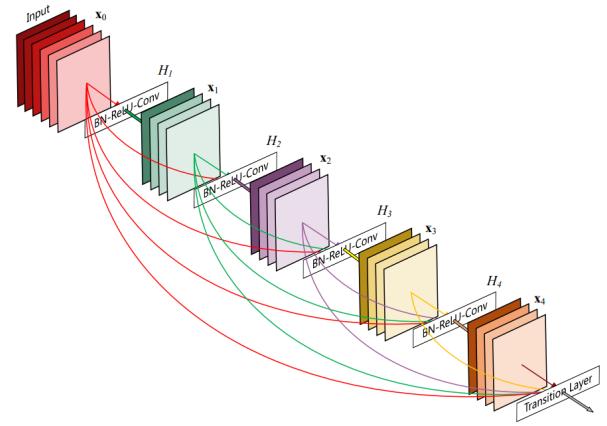


Figure 4: Illustration of DenseNet architecture [4], showcasing the innovative concept of dense connectivity. Each of the five depicted convolutional layers receives concatenated feature maps from all preceding layers, resulting in highly enriched feature representations. These layers are preceded by batch normalization and ReLU activation, with their outputs concatenated to form the subsequent layers' inputs, thus embodying the essence of dense connectivity.

and then ReLU is used for nonlinearity. Then, the output of the ReLU layer is convolved with a 1x1 kernel. The resulting output of the convolution layer with 1x1 kernel has $4k$ output channels. Then, the composite function defined above is applied to this convolution layer's output, resulting in k output channels. The bottleneck layer consists of batch normalization, ReLU, 1x1 convolution, and the composite function. DenseNets containing bottleneck layers are called "DenseNet-B."

9) Compression: To reduce the size of the model, a compression factor (θ) can be applied in the transition layers to reduce the number of feature maps:

$$m \rightarrow \text{transitionlayer} \rightarrow \lfloor \theta m \rfloor \quad (7)$$

where m is the number of input feature maps to the transition layer, and $\lfloor \cdot \rfloor$ denotes the floor function, ensuring an integer number of output feature maps.

DenseNets containing this compression operation are called "DenseNet-C," if the network both contains the bottleneck layer and the compression in the transition layers, then the network is called "DenseNet-BC."

10) Implementation: Implementing DenseNet involves stacking multiple densely connected blocks, each followed by a transition layer (except for the last block). Each block consists of multiple layers as defined by the composite function, and the entire network is capped by a global average pooling layer followed by a fully connected layer for classification. This modular and highly parameter-efficient architecture has been shown to achieve excellent performance on image classification tasks while being more computationally efficient than many traditional convolutional networks.

E. Inception [5]

Inception, also known as GoogleNet, is a convolutional neural network architecture developed by Christian Szegedy et al. and is introduced in the paper "Going deeper with convolutions" [5]. It was designed to participate in the ImageNet Large Scale Visual Recognition Challenge 2014 (ILSVRC 2014) and won the competition.

It had several improvements over the initial version that won the ILSVRC 2014, namely Inception-v3, Inception-v4, and Inception-Resnets. The v2 is briefly discussed and directly improved to v3 in the same paper, "Rethinking the Inception Architecture for Computer Vision" [6]. The later improvements are introduced in the paper "Inception-v4, Inception-ResNet and the Impact of Residual Connections on Learning" [7].

Each improvement of the Inception architecture increases its parameters, starting from 5.491M in the Inception-v1 to 21.640M in the Inception v3, and later to 43M in the inception v4. Inception-Resnets have a similar number of parameters to the v3 and v4. Hence, due to increasing computational demand, only Inception-v1 will be run and discussed in this project.

The detailed architecture of the model is provided in Table 8.

1) *Convolution Layers*: All the convolution layers in the model, whether in an inception module or stand-alone, are a sequence of a convolution, a batch normalization, and a ReLU.

2) *Inception Layers*: The key idea behind the Inception model is to have multiple different-sized convolutional filters operate on the same level and then concatenate their outputs as different channels. The resulting network gets *wider* instead of *deeper*. The Inception module introduced in the paper is given in Fig. 5.

It is noted in the paper that embeddings, as in dimension reductions and projections, represent information in a dense, compressed form, and compressed information is harder to model. However, if done judiciously, they can be done without much or any loss in representational power. This introduces the idea of performing 1x1 convolutions to reduce channels before the expensive 3x3 and 5x5 convolutions, significantly reducing the total computational load without sacrificing representativity. Hence, the improved Inception module is shown in Fig. 6, beside the *naive* implementation.

3) *Auxiliary Branches*: In the original paper, an idea to address the vanishing gradient problem is to add *auxiliary branches* to the network. These are separate outputs from the network branched from earlier layers with their ending sequence, that is, pooling and a fully connected network. The outputs from these auxiliary branches are also subjected to the same loss function and added to the total loss with a discount weight:

$$\text{TotalLoss} = \text{Loss}_{\text{main}} + 0.3 \cdot \text{Loss}_{\text{Aux}_1} + 0.3 \cdot \text{Loss}_{\text{Aux}_2} \quad (8)$$

However, these auxiliary branches are only considered during training and are discarded during validation and testing. The later published paper [6], that brings the inception-v3, retrospectively points out that the auxiliary branches act as a regularizer to the network.

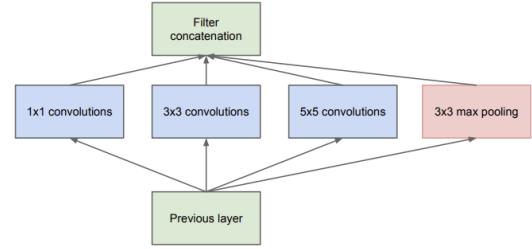


Figure 5: Inception Module, naive version.

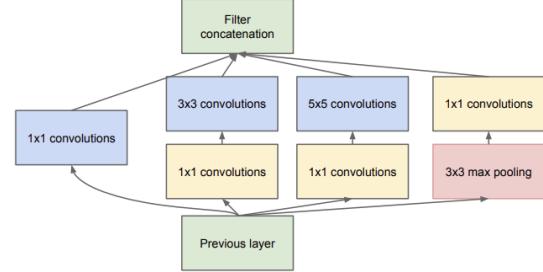


Figure 6: Improved Inception Module with the dimension reductions.

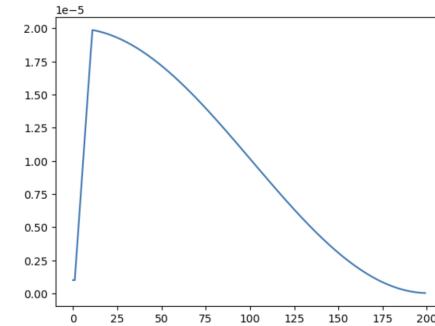


Figure 7: An example cosine annealing with a warm-up of 20 epochs.

4) *Label Smoothing*: The second paper introduces label smoothing to the Inception-v3 model. Label smoothing is a regularization method that introduces noise to the labels. The idea is to prevent the network from becoming overconfident in its predictions and reduce overfitting. It is done by replacing the hard 0 and 1 classification labels by $\frac{\epsilon}{K-1}$ and $1 - \epsilon$, respectively.

5) *Cosine Annealing with Warm-up*: Cosine Annealing with Warm-up is a learning-rate scheduling technique in which the learning rate is started with a near-zero value and linearly incremented to its base value during the warm-up epochs. Then, the learning rate is determined by the values of cosine for the rest of the epochs such that the rest of the epochs will be a quarter of a period of a cosine, as shown in Fig. 7.

III. RESULTS

In this section, we present a comprehensive analysis of the performance of various convolutional neural network (CNN) architectures, specifically VGGNet, ResNet, DenseNet, and

Model	Parameter Size	Train Acc. (%)	Val. Acc. (%)	Test Acc. (%)	Epoch No.	Dropout	Learning Rate	Momentum	Model Version	Pooling Method
1	22,205,450	97.86	83.96	84.01	29	0.5	0.01	-	VGG11	Max Pooling
2	22,205,450	98.81	80.8	80.3	23	0	0.01	-	VGG11	Max Pooling
3	22,205,450	94.58	83.94	83.74	20	0.5	0.1	-	VGG11	Max Pooling
4	22,205,450	99.86	86.45	86.2	21	0	0.1	-	VGG11	Max Pooling
5	22,205,450	98.19	79.52	79.25	17	0	0.1	0.9	VGG11	Max Pooling
6	22,205,450	99.44	85.38	86.16	16	0	0.1	-	VGG11	Average Pooling
7	33,014,858	99.29	85.48	84.3	19	0	0.1	-	VGG19	Max Pooling

Table 3: Comparison of VGGNet Models

Model	Parameter Size	Train Acc. (%)	Val. Acc. (%)	Test Acc. (%)	Epoch No.	Dropout	Learning Rate	Momentum	Model Version	Pooling Method
1	11,173,962	94.35%	88.58%	89.06%	25	0	0.001	0	Resnet18	Avg. Pooling
2	11,173,962	94.73%	88.48%	88.50%	25	0	0.01	0	Resnet18	Avg. Pooling
3	11,173,962	92.96%	88.54%	88.80%	25	0.5	0.01	0	Resnet18	Avg. Pooling
4	11,173,962	94.95%	88.96%	89.47%	25	0.5	0.001	0	Resnet18	Avg. Pooling
5	11,173,962	92.72%	87.16%	87.29%	25	0.5	0.001	0.9	Resnet18	Avg. Pooling
6	11,173,962	73.06%	59.24%	59.01%	25	0.5	0.001	0	Resnet18	Max Pooling
7	21,282,122	94.76%	88.98%	89.23%	25	0.5	0.001	0	Resnet34	Avg. Pooling

Table 4: Comparison of ResNet Models

Model	Parameter Size	Train Acc. (%)	Val. Acc. (%)	Test Acc. (%)	Epoch No.	Dropout	Learning Rate	Momentum	Model Depth	Model Version
1	4,113,082	89.62	81.62	82.14	18	0.5	0.01	-	L=36, k=24	DenseNet
2	4,113,082	84.8	78.2	78.34	16	0	0.01	-	L=36, k=24	DenseNet
3	4,113,082	92.47	82.64	82.71	24	0.5	0.1	-	L=36, k=24	DenseNet
4	4,113,082	97.24	85.24	85.78	25	0	0.1	-	L=36, k=24	DenseNet
5	4,113,082	94.34	84.3	83.71	27	0	0.1	0.9	L=36, k=24	DenseNet
6	1,979,098	94.26	83.32	84.38	17	0	0.1	-	L=36, k=24	DenseNet-BC
7	9,016,714	88.5	82.9	82.45	15	0	0.1	-	L=54, k=24	DenseNet

Table 5: Comparison of DenseNet Models

Model	Parameter Size	Train Acc. (%)	Val. Acc. (%)	Test Acc. (%)	Epoch No.	Dropout	Learning Rate	Momentum	Model Depth	Model Version
1	5,491,000	99.0785	92.7684	92.8385	50		No			No
2	5,491,000	99.5072	92.0974	92.0974	50		No			Yes
3	5,491,000	98.5757	92.3177	92.3778	50		Yes			No
4	5,491,000	99.1647	91.9671	91.9772	50		Yes			Yes

Table 6: Comparison of Inception Models

Inception models, tailored and evaluated on the CIFAR-10 dataset. The experimental setup was designed to explore the effects of different hyperparameters and architectural choices—such as depth, dropout, learning rate, and pooling methods—on the models’ training and generalization capabilities. While the experiments were not optimized to achieve peak performance, they provide valuable insights into the models’ behavior under various configurations.

A. VGGNet Results

The Table 3 presents a comparative analysis of different VGGNet models on the CIFAR-10 dataset, evaluating their performance in terms of parameter size, accuracy, and training configuration. Models 1 to 7 vary in terms of dropout rate, learning rate, and the use of momentum during training. Models 1 and 2, which employ a lower learning rate of 0.01, show a significant difference in performance when dropout is introduced. Model 1, with a dropout rate of 0.5, achieves a training accuracy of 97.86% and a test accuracy of 84.01%. Model 2, which has no dropout, sees a slight increase in training accuracy to 98.81% but a noticeable decrease in test accuracy to 80.3%. Models 3, 4, and 5, with a higher learning rate of 0.1, display a diverse impact of dropout and momentum. Model 3, with dropout, shows lower training accuracy (94.58%) but maintains a competitive test accuracy (83.74%). Model 4, without dropout, demonstrates the highest training accuracy (99.86%) and test accuracy (86.2%) among the VGG11 models

with max pooling. Model 5, which introduces momentum (0.9) experiences a decrease in test accuracy to 79.25%, suggesting that momentum may not contribute positively to generalization for this model configuration. Model 6, employing an average pooling method instead of max pooling, achieves a slightly lower test accuracy (86.16%) compared to the best-performing VGG11 model with max pooling (Model 4). Model 7, which expands the architecture complexity to VGG19, maintains a relatively high training accuracy (99.29%) but exhibits a lower test accuracy (84.3%) compared to the best VGG11 model (Model 4). This indicates that increasing the model depth to 19 layers does not necessarily translate into better performance. All models were trained by using early stopping and they have stopped between 15-20 epochs. All these models were trained using Google Colab with V100 GPU.

B. ResNet Results

Table 4 displays the results of 7 different ResNet models trained with different parameters.

The training of the ResNet-18 models was done using the Adam (short for Adaptive Moment Estimation) optimization algorithm, which is commonly used for training deep neural networks. Two values were tested for learning rate, momentum constant, and dropout probability. Average and max pooling were compared, keeping the other parameters at their best-performing values. The resulting best-performing model parameters were used to train the ResNet-34 model.

A lower learning rate (0.001) has achieved slightly better results than greater learning rates (0.01). A dropout probability of 0.5 has performed better compared to no dropout case. The best of the four cases was again trained with momentum using 0.9 as the constant. The test accuracy in this case has dropped around 2%. Max pooling was tested using a learning rate of 0.001, a dropout probability of 0.5, and no momentum. The resulting accuracy has dropped dramatically compared to previous cases, around 25 to 30%.

When the best-performing parameters were used in training ResNet-34 architecture, it did not result in a big change in the accuracy rate of the model. It was at the same percentile with the best-performing ResNet-18 model, at 89%.

All models were trained for 25 epochs. The models were trained using Google Colab with T4 GPU.

C. DenseNet Results

Table 5 is the training results of seven different DenseNet-based models with slight differences. The experiment's goal was not to obtain the best results, therefore, the parameter selection was made for experimental purposes. Additionally, a dropout layer was included before the classification layer of the DenseNet, which does not exist originally in [4]. Additionally, the parameter L represents the depth of the network. The depth represents the number of DenseBlocks, which is Batch Normalization, ReLU and 3x3 convolution layer, in this experiment. The transition layers are not included in calculating the parameter. The depth count standard is not given in [4]. Therefore, this parameter is interpreted as the number of DenseBlocks. [4] states that they experimented with CIFAR-10 but did not share their models' exact architecture. Their primary focus was ImageNet. Therefore, the CIFAR-10 experiment is not given in detail. They said they used a network with $L = 40$. However, as said before, the standardization of the parameter L is not clearly defined. Therefore, the network used in this experiment is constructed somewhat instinctively. Each model in Table 5 has three DenseLayers, which consist of DenseLayers and use input concatenation at the output after each block. Transition layers determine the limits of the DenseLayers. There are 6 DenseBlocks in models 1-6 and 12 in Model 7. In the second DenseLayer, there are 18 DenseBlocks in models 1-6 and 24 in model 7. Finally, there are 12 DenseBlocks in the third DenseLayer in model 1-6, and 24 DenseBlocks in the third DenseLayer of model 7. The motivation behind this is that the first three DenseLayers in the implementation of DenseNet for ImageNet were copied, omitting the fourth DenseLayer considering that CIFAR-10 is much simpler than the ImageNet and, therefore, the first three layers should be enough to capture the relations between the images. The intuition behind model 7 is that the more DenseBlocks should capture more features, leading to better results. All models except model 6 are in the Default mode, meaning they do not utilize the bottleneck layer and the compression in the transition layer. Model 6 was designed to test the efficiency of the bottleneck layer and the compression in the transition layer. However, testing these in different models might be more convenient to see the individual effects.

Due to strict time constraints, the experiments on more models could not be performed. The expansion rate k selection is directly from [4]. Again, experimenting with different k values might be more beneficial to see the effect on the performance of the image classification task. Other parameters are chosen according to the experiments discussed above so that the performance of the models can be compared easily.

Google Colab was utilized to train these models due to its high-performance T4 GPUs. The training parameters are chosen similarly to the implementation of the VGGNet. The batch size 64 and the SGD optimizer from the "PyTorch.optimizer" module were used. Each model was run for 60 epochs, but early stopping based on validation loss is used. The patience was set 7, meaning that the training was stopped if the validation loss did not improve for 7 epochs. Then, the best performing model in terms of validation loss was outputted at the end of the training using the PyTorch's "load_state_dict()" function. Table 5 gives the learning rate and momentum parameters. The first four models for tests to find the best learning rate and dropout rate. Then, the best-performing model, model 4, was tested with momentum in model 5. Model 4 and Model 5 were compared and Model 4 was found to be superior. Again, Model 4 parameters in the DenseNet-BC configuration were tested in Model 6. Model 4 was again found to be superior. Then, the number of DenseBlocks in the DenseLayers increased in Model 7.

Comparing Model 1 (82.14%) and Model 2 (78.34%), the model with dropout probability performed better. However, Model 4 (85.78%) performed better comparing to Model 3 (82.71%) even though Model 4 has no dropout. Comparing learning rates, the ones with the higher learning rates performed better. Model 5 (83.71%) was trained with momentum, but this does not result in an improvement in the test accuracy. DenseNet-BC setup was experimented in Model 6 (84.38%), performed even though it has a 1.9M parameters, which is comparatively very low. Model 7 (82.45%) contains more parameters comparing the to the rest of the DenseNet models, however, it does not perform as expected.

D. Inception Results

The training was done using Stochastic Gradient Descent with 0.9 momentum. Cosine annealing with a warm-up method was used for scheduling with 30% of the total epochs used for warm-up and 0.01 as the base value of the learning rate. Nvidia GeForce RTX 4090 was used as the training GPU.

Metrics such as loss and accuracy, and confusion matrices are provided in Fig. 15-22. A simplified results of the four model variations are provided in Table 6. The more extensive results are provided in Table 7.

All the model variations had relatively close test set accuracy results that no two result is farther away than %1. Best test set accuracy was obtained by the vanilla Inception model by %92.8385. Vanilla Inception is also the best performer in all Top-k accuracy classifications. The second model is with label smoothing but without the auxiliary branches with %92.3778. However, this model variation performed second only in the Top-1 and Top-2 accuracy, and is in the last place for the Top-3

accuracy. The model with auxiliary branches but without label smoothing is in the third place, and the model with both are in the third and fourth place respectively. Their other Top-k accuracy metrics are also in the same fashion, expect the situation with the second model.

All the model variations had pretty similar graphs for their loss and accuracy per epoch. However, it must be noted that direct comparisons of loss values would not be correct since both the auxiliary branches and label smoothing introduce additional loss values into the calculation, and the same model performance otherwise would have different loss values.

IV. DISCUSSION

This experiment was run to classify the CIFAR-10 dataset images into ten classes. The dataset contains 60000 32x32 RGB images. The pre-processing for the data was not deemed necessary since one image has $32 \times 32 \times 3 = 3072$ features. This feature size for a sample seems to be enough for the classification task, considering there are only 10 classes for 60000 images. Convolutional networks are relied upon for their efficacy in accurately classifying images. Therefore, no new feature engineering or analysis is run on the dataset in terms of pre-processing the image. Additionally, the methods used in this experiment are not run on other datasets. However, popular datasets like ImageNet, CIFAR-100, and SVHN were tested in the descendants of the methods implemented in this paper. Especially, ImageNet was tested on various models due to the famous ImageNet competition held annually.

In the study of VGGNet architecture, variations of the learning rate, dropout probability, momentum, and pooling methods reveal different outcomes where each hyperparameter significantly impacts model performance. Adam optimization algorithm has been chosen to optimize the performance of the network. A higher learning rate of 0.1 has shown that a higher learning rate has performed better than 0.01 by reducing training error and boosting test accuracy when coupled with dropout probability 0. A high learning rate, allows higher steps towards convergence and makes it easier to escape from the local minimum. Then, the introduction of the dropout probability 0.5 for these learning rate scenarios has shown that a smaller learning rate 0.01 was better when dropout is introduced compared to 0.1. It allows smaller weight updates and smaller steps towards convergence to avoid overshooting. Dropout regularization compels the neurons with a specific probability to prevent overfitting. For the model, the learning rate is 0.1, and the dropout probability is varied; eliminating half of the neurons leads to a significant loss of information and learning capacity. For the models whose learning rate is 0.01, the introduction of dropout regularization increased the performance of the network by preventing overfitting and increasing the network's generalization ability. Then, momentum was coupled with a high learning rate and 0 dropout probability. The performance of the network has decreased. Introducing the momentum allows weight updates depending on both current and previous updates. For this particular case, momentum has caused the network to reach optimal points during training, leading to less accurate results. Then, the best

model was used by changing the max-pooling layers with average pooling, which typically calculates the average instead of taking the maximum value. For this particular model, changing the max pooling layer with average pooling layers remained stable with a slight decrease compared to max pooling. The average pooling layer might cause to forget some critical features that max pooling would capture aggressively. Then, the best model is trained with an expanded network which contains 16 convolution layers rather than 8. This modification has increased the required training time significantly without improving the performance. Adding additional weights to the network might cause being more inclined to overfitting, especially without dropout. This experience underlines the importance of selecting hyperparameters carefully by considering the network.

Experimenting with different configurations of ResNet-18 and ResNet-34 has yielded useful information for comparing different parameters and their effect on the model performance. The choice of learning rate, dropout probability, momentum, and pooling method have influenced the training outcomes. Adam optimization algorithm was chosen and was not replaced as commonly used in CNN models. The variation in learning rates has shown that a lower learning rate of 0.001 has performed slightly worse compared to a higher learning rate of 0.01, which is in line with the expectation as a lower learning rate allows smaller weight updates and smaller steps towards convergence to avoid overshooting. Introducing dropout regularization has proved to improve the performance across most cases compared to scenarios with no dropout. This result is in line with the expectation, as the dropout's role is to prevent overfitting by randomly dropping neurons during training, increasing the network's generalization ability. A slight decrease in accuracy was observed when a lower learning rate and dropout regularization were coupled with momentum. The effect of momentum on gradient descent can be equivalent to higher learning rates, which decreases the model's accuracy rate. When max-pooling was used instead of average pooling after the last block in the ResNet architecture, the accuracy rates dropped dramatically, implying that average pooling preserves more spatial information about the input images compared to max pooling, which results in better overall feature extraction. When the optimized parameters were carried into ResNet-34 architecture, no significant change in the accuracy rate was observed, indicating that the optimization of the model may be saturated after some level of model complexity. Improving the results further might require exploring architectural adjustments or different model architectures.

An analysis was made on the CIFAR-10 dataset using the DenseNet architecture with slight differences. The architecture was offered to capture more insight into the dataset using fewer parameters with skip connections. However, it utilizes skip connections slightly differently from the infamous ResNets. The DenseNet concatenates its inputs to the outputs in some layers so that the information about the input will be used in the later layers of the network. This also promises to solve the vanishing gradient problem in Deep Networks. In addition to the original network architecture offered in [4], a dropout layer

was introduced before the classification layer of the network. This change was made with the intention of preventing overfitting in the complex network structures. As the complexity increases in the network, it has more parameters to memorize the train set. The network also has bottleneck connections, meaning that the features are firstly increased so that the later layer has more data to capture the relations. The network complexity in terms of layer sizes was not clearly specified in the original implementation of the DenseNet; therefore, different size parameters were defined in this implementation. The network also offers a compression method that reduces the parameter size significantly. For example, Model 6 utilizes a compression rate of 0.5 and uses bottleneck connections. However, the number of parameters of Model 6 is found to be 1.9 million. However, Model 4 has the same depth but does not utilize compression. Its parameter size is around 4 million. Also, looking at the test accuracy scores, it can be said that the compression serves the purpose since Model 6 performed similarly to Model 4, even though Model 6 has fewer parameters. The learning differed in the analysis. Model 1 and Model 2 reached early stopping before Model 3 and Model 4, probably because a higher learning rate resulted in more oscillatory behaviors of the weight updates, resulting in a more unstable training process. However, the ones with higher learning rates converged to local minima with better testing accuracy. Another parameter tested is the dropout rate, which originally did not exist in the DenseNet implementation. The effect of dropout is not consistent. Model 1 is superior to Model 2, but Model 3 is inferior to Model 4. The increase might be because the dropout prevents overfitting and results in better test accuracy, or it might be stochasticity since the initialization is from random distributions. The decrease might be because dropouts might prevent learning due to dropped connections in the training. The momentum was tested to achieve better performance in Model 5, but it could not surpass the performance of Model 4. The momentum constant was chosen as 0.9. This might dominate the gradient, resulting in the weight update always going in the same direction, even though the gradient might indicate another direction. A new block structure and compression were tested with Model 6. The parameter size decreased to half. However, the network performance remained unchanged. The model might perform better if the parameter size were set to 4 million. The last model was tested to see the effect of network size. Model 7 has two times the parameters of Model 4. The performance of the model did not increase as expected. Overall, the best performance was obtained from Model 4 with 85.78%. However, in the original implementation, reached error rates of 5.77% with default layers and reached 5.19 % suggesting that this experiment needs further hyperparameter tuning. In addition, the original implementation used 7 million and 15 million parameters, showcasing the effectiveness of the network size. However, these results were obtained when the training was done for 300 epochs, which is not applicable to this experiment due to time limits. Similar results might be obtained if the early stopping was removed and the training was done for 300 epochs.

Experiments with the various Inception models yielded quite

unexpected results. Surprisingly, the vanilla setup obtained the best test set accuracy in all accuracy metrics, as shown in Table 7. This is also in contrast with what the original papers suggest since in both [5] and [6], it is argued that both auxiliary branches and label smoothing improve the results. Auxiliary Branches are introduced in the paper [5]; however, they are further discussed in the paper [6]. It is argued in [6] that auxiliary branches have a regulatory role in the model, and their effects only start to be noticeable rather late in the training. It is 100 epochs of training in the original papers in contrast to 50 in my experimentation. Hence, their apparent ineffectiveness might be caused by the lack of late-training epochs in which auxiliary branches perform their role as regularizers and provide better accuracy by reducing overfitting.

Label smoothing also performed ineffectively in contrast to the expectation. It is quite surprising since label smoothing is a somewhat popular method that is known to prevent the network from becoming over-confident to improve its results, and it is used in many state-of-the-art models, including image recognition, language translation, and speech recognition. [8] A similar argument about the lack of late-training epochs can be made for label smoothing since its main role in the model is also as a regularizer.

The vanilla Inception achieved the best test set accuracy of %92.8385 among all models in the project.

In a comparative study of CNN architectures on the CIFAR-10 dataset, the analysis revealed useful conclusions where model complexity, such as that seen in VGG19, did not always correlate with superior performance, as simpler VGG11 configurations often yielded better results. DenseNet models stood out for their efficient parameter usage and robustness, while ResNet models benefited from residual connections but required careful consideration of pooling methods and momentum for optimal generalization. Inception models indicated that additional complexities like auxiliary branches and label smoothing might not enhance performance within certain training limits, as the vanilla Inception setup achieved the highest accuracy. This suggests that for specific datasets like CIFAR-10, simpler architectures or those with thoughtful implementation of complexity, can be more effective, emphasizing the importance of matching the model design intricately with the dataset characteristics.

REFERENCES

- [1] *The CIFAR-10 dataset*, Alex Krizhevsky, Vinod Nair, and Geoffrey Hinton, University of Toronto, 2009. [Online]. Available: <https://www.cs.toronto.edu/~kriz/cifar.html>
- [2] K. Simonyan and A. Zisserman, "Very Deep Convolutional Networks for Large-Scale Image Recognition." arXiv, Apr. 10, 2015. doi: 10.48550/arXiv.1409.1556.
- [3] K. He, X. Zhang, S. Ren and J. Sun, "Deep Residual Learning for Image Recognition," 2016 IEEE Conference on Computer Vision and Pattern Recognition (CVPR), Las Vegas, NV, USA, 2016, pp. 770-778, doi: 10.1109/CVPR.2016.90.
- [4] G. Huang, Z. Liu, L. Van Der Maaten and K. Q. Weinberger, "Densely Connected Convolutional Networks," 2017 IEEE Conference on Computer Vision and Pattern Recognition (CVPR), Honolulu, HI, USA, 2017, pp. 2261-2269, doi: 10.1109/CVPR.2017.243.
- [5] C. Szegedy et al., "Going deeper with convolutions," 2015 IEEE Conference on Computer Vision and Pattern Recognition (CVPR), Boston, MA, USA, 2015, pp. 1-9, doi: 10.1109/CVPR.2015.7298594.
- [6] C. Szegedy, V. Vanhoucke, S. Ioffe, J. Shlens and Z. Wojna, "Rethinking the Inception Architecture for Computer Vision," 2016 IEEE Conference on Computer Vision and Pattern Recognition (CVPR), Las Vegas, NV, USA, 2016, pp. 2818-2826, doi: 10.1109/CVPR.2016.308.
- [7] Christian Szegedy, Sergey Ioffe, Vincent Vanhoucke, and Alexander A. Alemi. 2017. Inception-v4, Inception-ResNet and the Impact of Residual Connections on Learning. In Proceedings of the Thirty-First AAAI Conference on Artificial Intelligence (San Francisco, California, USA) (AAAI'17). AAAI Press, 4278–4284.
- [8] Rafael Müller, Simon Kornblith, and Geoffrey Hinton. When does label smoothing help? In Proceedings of the 33rd International Conference on Neural Information Processing Systems, 2019.

APPENDIX

A. Elaborated Results

This part gives detailed results on the microscopic level. The confusion matrices, the training losses, the validation losses, the training accuracy, and the validation accuracy are given in this part with the related model's section.

1) **VGGNet:** This section gives the results of the VGGNet. Figures 8-14 are the outcomes of the training of 7 different VGGNet models.

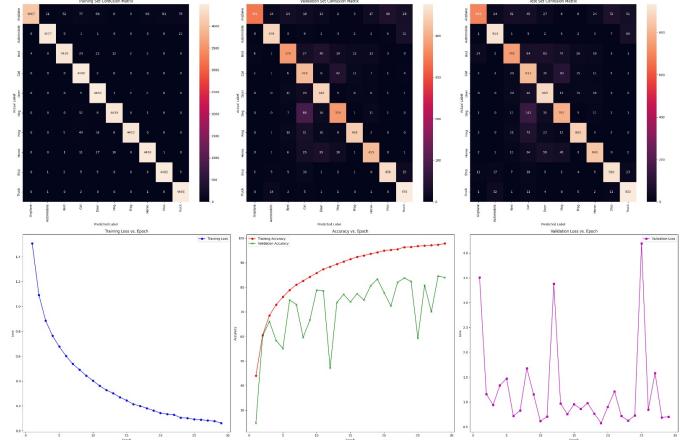


Figure 8: VGGNet Model 1 training results.

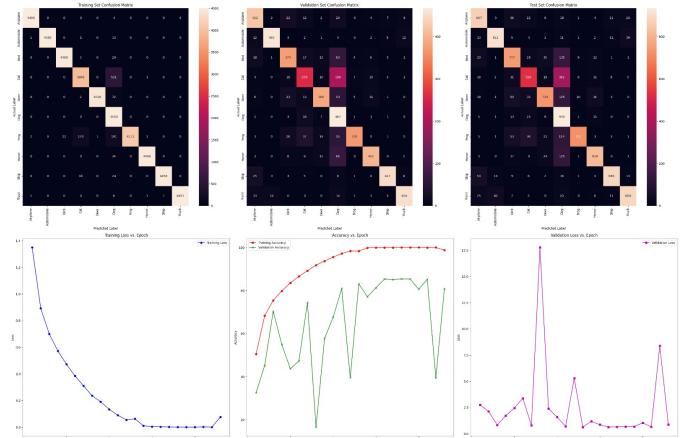


Figure 9: VGGNet Model 2 training results.

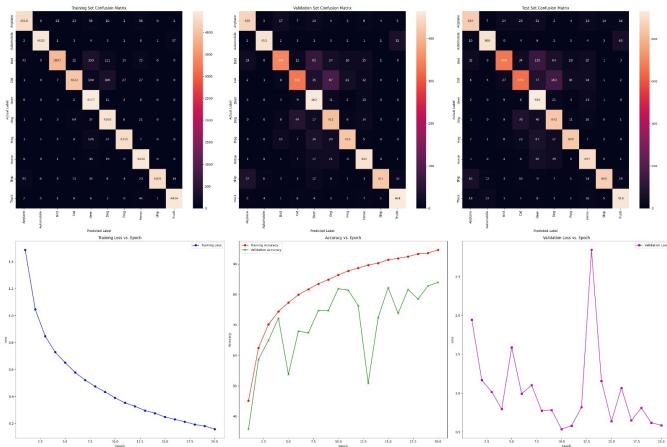


Figure 10: VGGNet Model 3 training results.

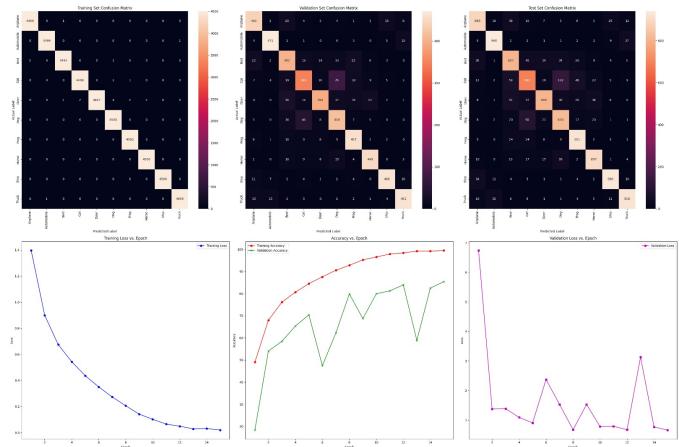


Figure 13: VGGNet Model 6 training results.

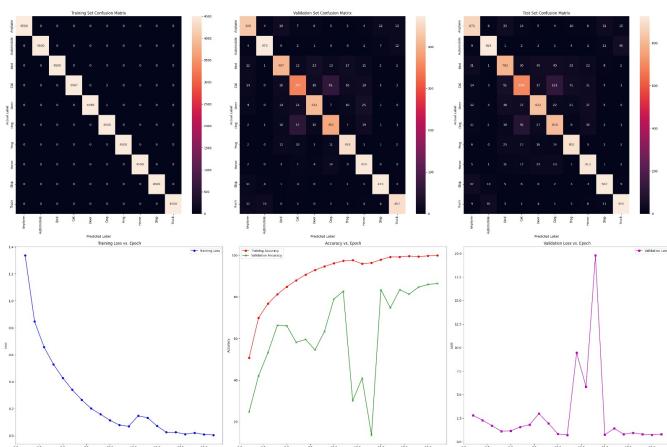


Figure 11: VGGNet Model 4 training results.

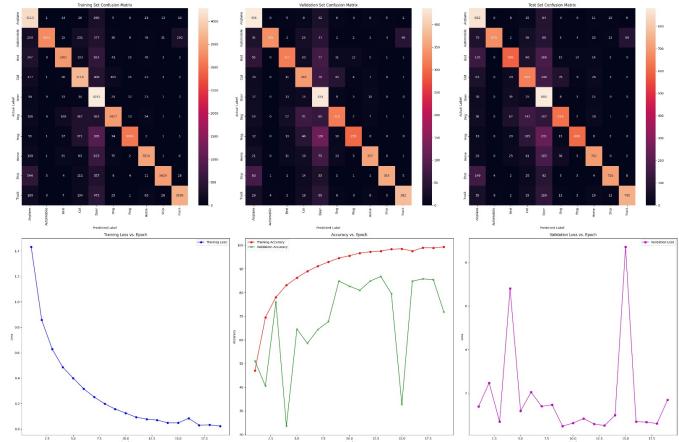


Figure 14: VGGNet Model 7 training results.

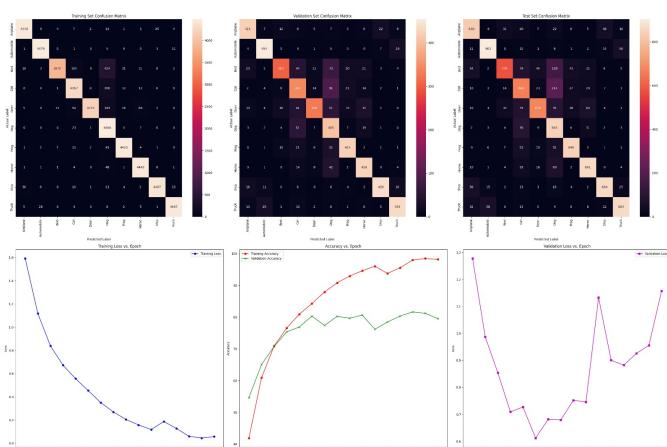


Figure 12: VGGNet Model 5 training results.

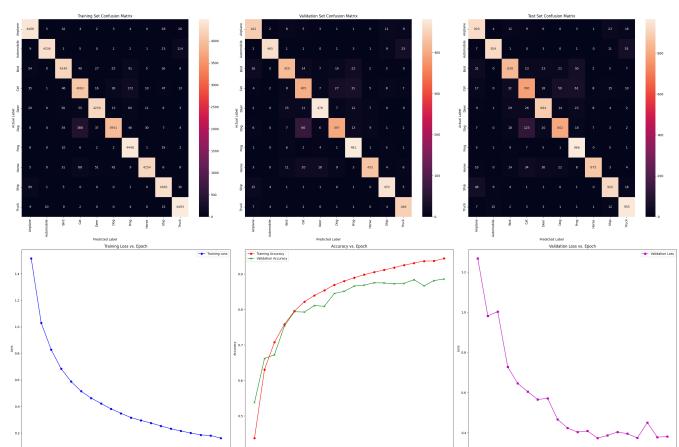


Figure 15: ResNet-18 Model 1 training results.

2) *ResNet*: This section gives the results of the ResNet. Figures 15-21 are the outcomes of the training of 7 different ResNet models.

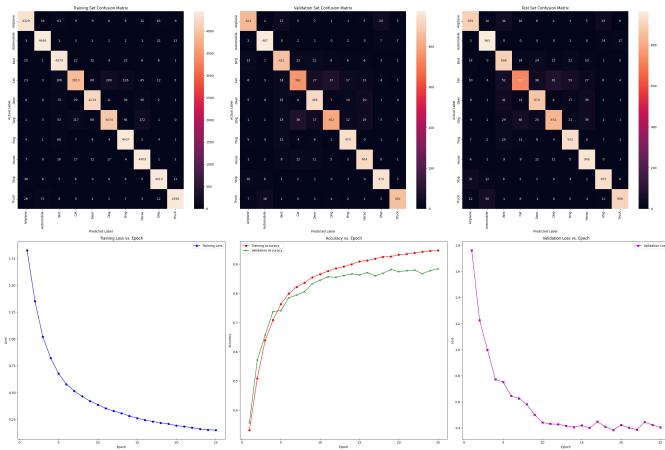


Figure 16: ResNet-18 Model 2 training results.

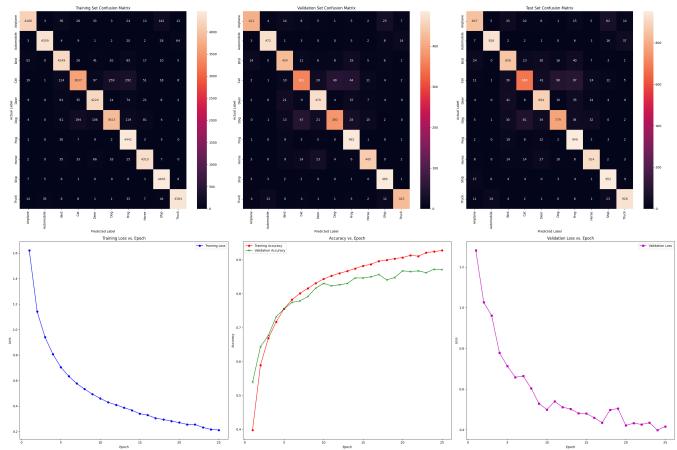


Figure 19: ResNet-18 Model 5 training results.

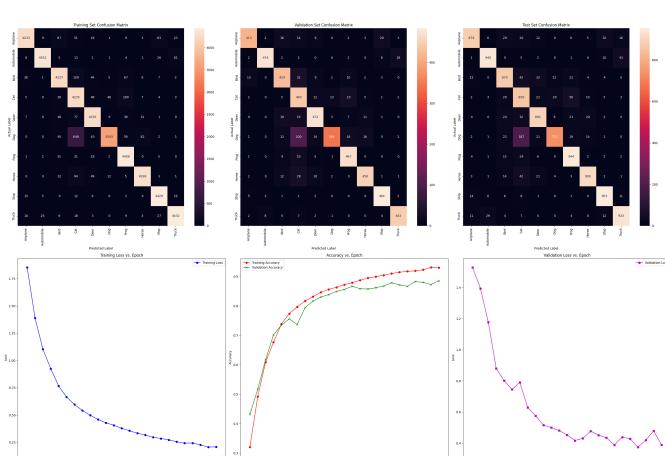


Figure 17: ResNet-18 Model 3 training results.

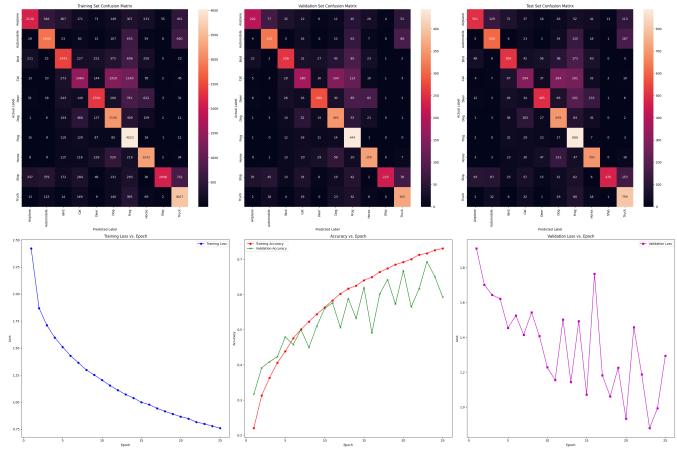


Figure 20: ResNet-18 Model 6 training results.

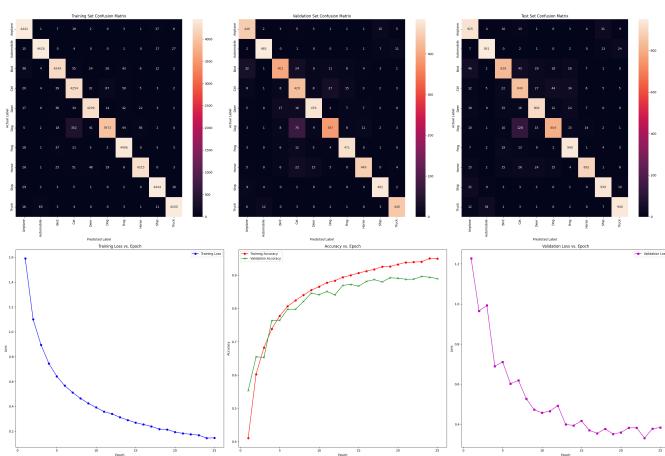


Figure 18: ResNet-18 Model 4 training results.

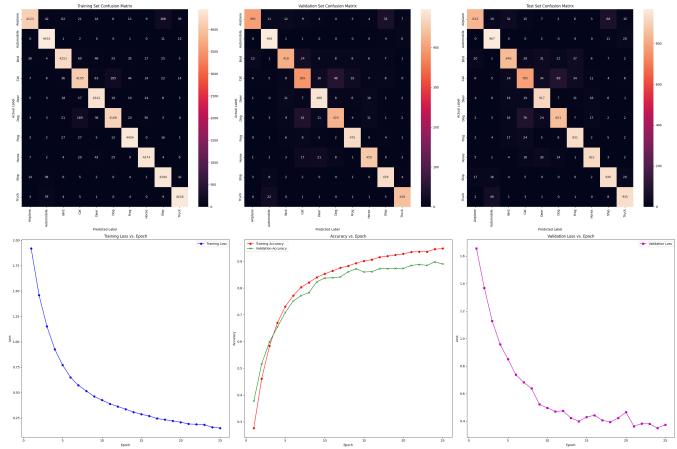


Figure 21: ResNet-34 Model 7 training results.

3) DenseNet: This section gives the results of the DenseNet. Figures 22-22 are the outcomes of the training of 7 different DenseNet models.

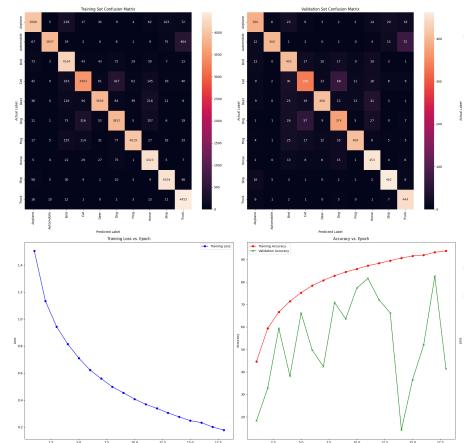


Figure 22: DenseNet Model 1 training results.

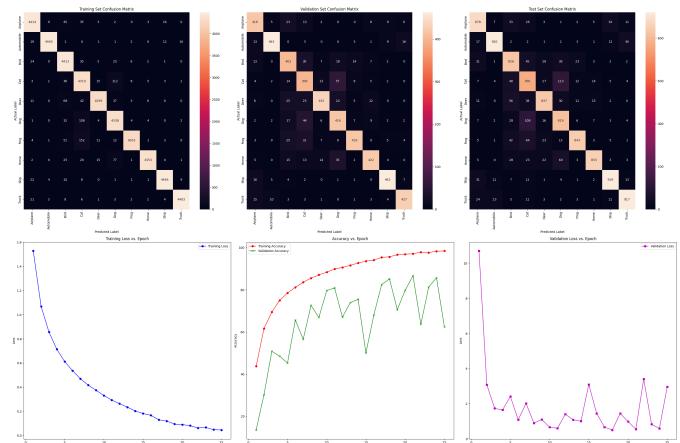


Figure 25: DenseNet Model 4 training results.

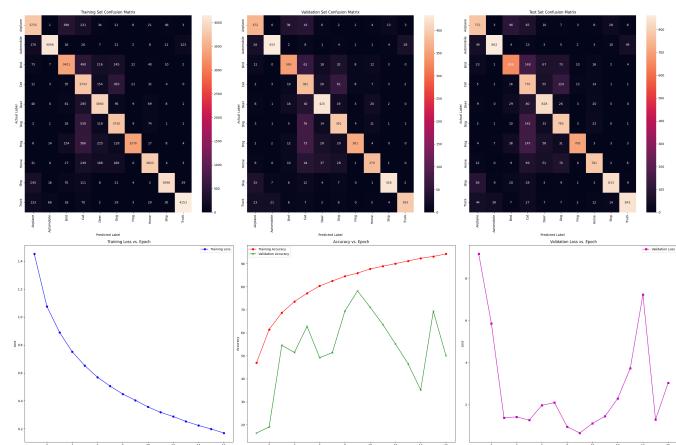


Figure 23: DenseNet Model 2 training results.

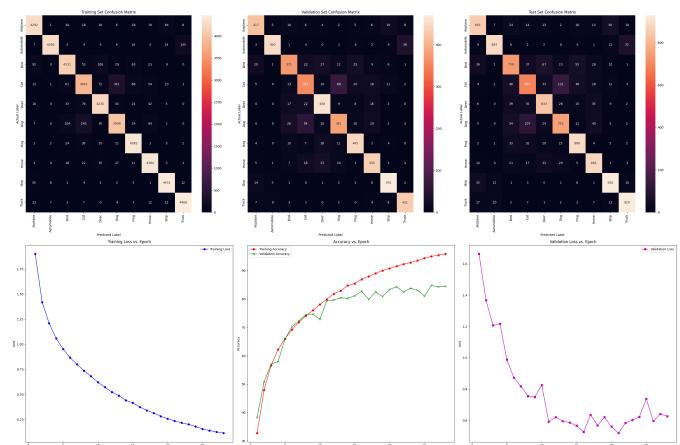


Figure 26: DenseNet Model 5 training results.

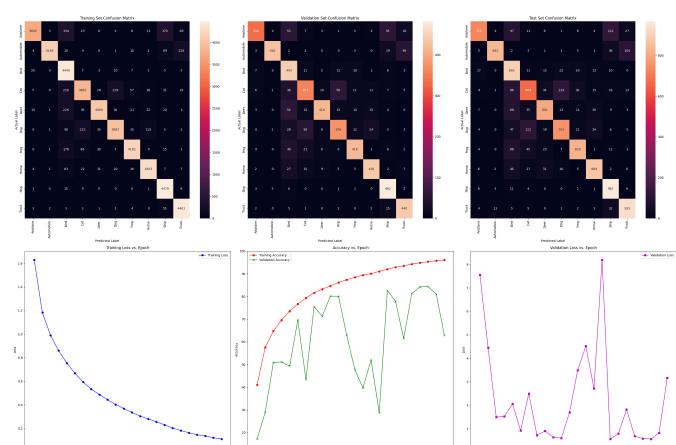


Figure 24: DenseNet Model 3 training results.

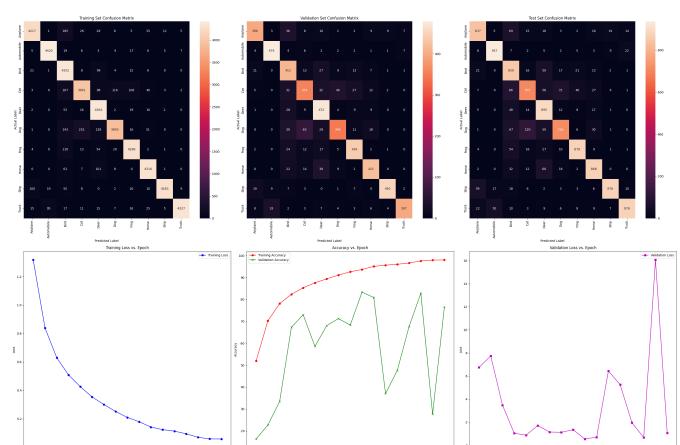


Figure 27: DenseNet Model 6 training results.

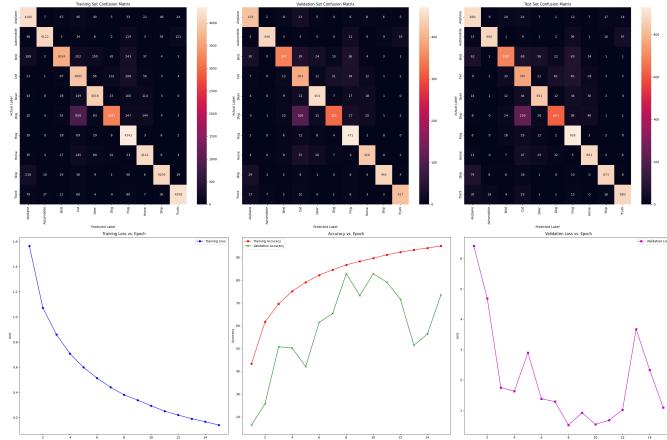


Figure 28: DenseNet Model 7 training results.

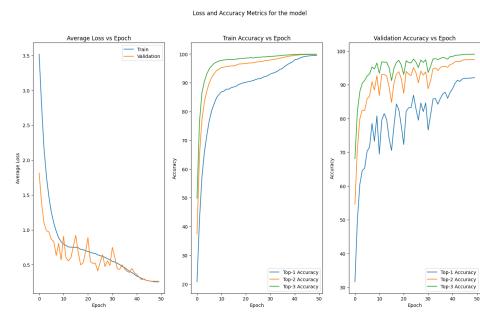


Figure 31: Inception Model 2 training results.

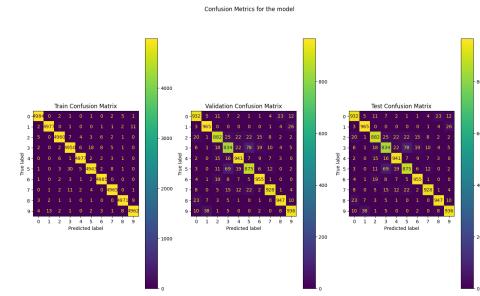


Figure 32: Inception Model 2 confusion matrices.

4) Inception: This section gives the results of the Inception. Figures 30-36 and Table 7-9 are the outcomes of the training of 4 different Inception models.

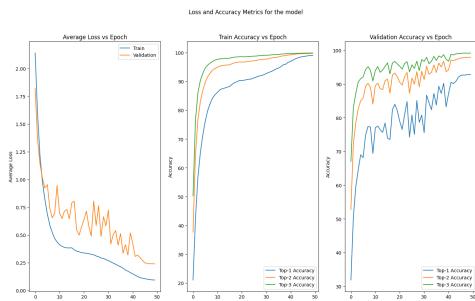


Figure 29: Inception Model 1 training results.

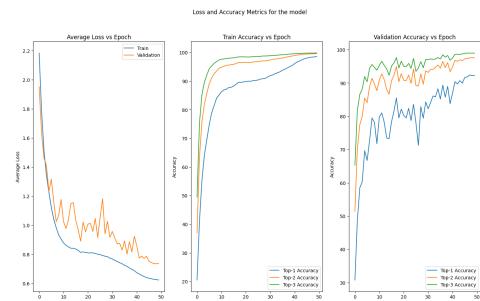


Figure 33: Inception Model 3 training results.

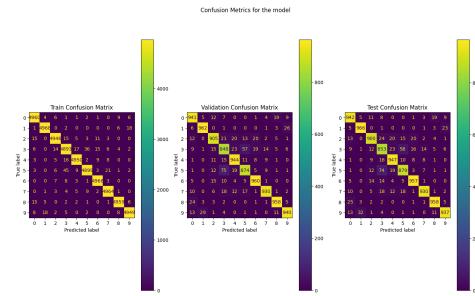


Figure 30: Inception Model 1 confusion matrices.

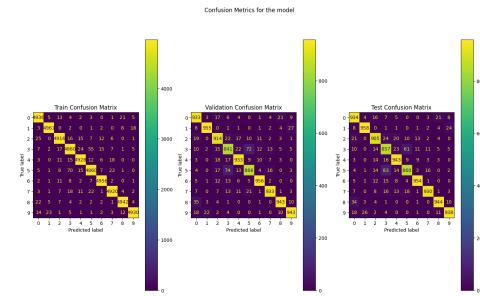
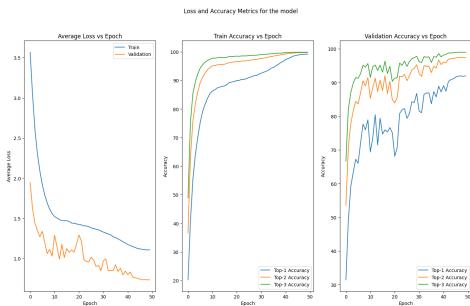
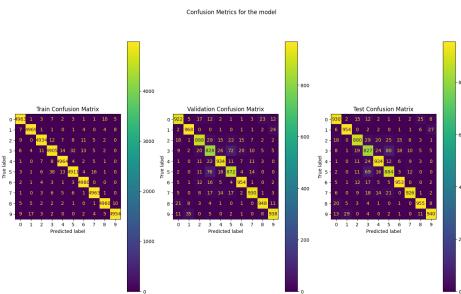


Figure 34: Inception Model 3 confusion matrices.

Model	Parameter Size	Train Top-1 Acc. (%)	Train Top-2 Acc. (%)	Train Top-3 Acc. (%)	Epoch No.	Label Smoothing	Auxiliary Branches
1	5,491,000	99.0785	99.8097	99.9079	50	No	No
2	5,491,000	99.5072	99.8898	99.9399	50	No	Yes
3	5,491,000	98.5757	99.5693	99.7937	50	Yes	No
4	5,491,000	99.1647	99.7496	99.8618	50	Yes	Yes
Model	Parameter Size	Val. Top-1 Acc. (%)	Val. Top-2 Acc. (%)	Val. Top-3 Acc. (%)	Epoch No.	Label Smoothing	Auxiliary Branches
1	5,491,000	92.7684	97.9367	99.1987	50	No	No
2	5,491,000	92.0974	97.4259	99.0385	50	No	Yes
3	5,491,000	92.3177	97.5361	98.9383	50	Yes	No
4	5,491,000	91.9671	97.3958	99.0084	50	Yes	Yes
Model	Parameter Size	Test Top-1 Acc. (%)	Test Top-2 Acc. (%)	Test Top-3 Acc. (%)	Epoch No.	Label Smoothing	Auxiliary Branches
1	5,491,000	92.8385	97.8966	99.1486	50	No	No
2	5,491,000	92.0974	97.4259	99.0385	50	No	Yes
3	5,491,000	92.3778	97.5361	98.9383	50	Yes	No
4	5,491,000	91.9772	97.4359	99.0184	50	Yes	Yes

Table 7: Detailed Results of Inception Models**Figure 35:** Inception Model 4 training results.**Figure 36:** Inception Model 4 confusion matrices.

B. VGGNet Python Implementation

```

1  # -*- coding: utf-8 -*-
2  """Untitled0.ipynb
3
4  Automatically generated by Colaboratory.
5
6  Original file is located at
7      https://colab.research.google.com/drive/197GaKSJ
     ↪ e1nU-7GIpSq3fu_t7UfMHlL_10
8  """
9
10 # Import necessary libraries for neural network
    ↪ construction, optimization, and dataset
    ↪ handling.
11 import torch
12 import torch.nn as nn
13 import torch.optim as optim
14 import torchvision
15 import torchvision.transforms as transforms
16 import matplotlib.pyplot as plt
17 import torch.nn.functional as F
18 import numpy as np
19
20 import torch
21 import torch.nn as nn
22 from torch.utils.data import DataLoader,
     ↪ random_split
23 import torch.optim as optim
24 from torchvision import datasets, transforms
25 from tqdm import tqdm
26 from torch.utils.data.sampler import
     ↪ SubsetRandomSampler
27 from torch.utils.data import Subset
28
29 # Define normalization transformation for the
     ↪ CIFAR10 dataset.
30 normalize = transforms.Normalize(
31     mean=[0.4914, 0.4822, 0.4465],
32     std=[0.2023, 0.1994, 0.2010],
33 )
34
35 # Compose the transformations: resizing the images,
     ↪ converting them to tensor, and normalizing.
36 transform = transforms.Compose([
37     transforms.Resize((227, 227)),
38     transforms.ToTensor(),
39     normalize,
40 ])
41
42 # Function to split the dataset by class, ensuring
     ↪ each class is represented in both the train and
     ↪ validation sets.
43 def split_by_class(dataset, val_percentage):
44     # Determine the number of classes and
     ↪ initialize indices for each class
45     class_indices = {i: [] for i in range(10)}
46
47     # Iterate through the dataset and assign
     ↪ indices to the respective class
48     for idx, (_, label) in enumerate(dataset):
49         class_indices[label].append(idx)
50
51     # Initialize train and validation indices
52     train_indices, val_indices = [], []
53
54     # Split the indices for each class into train
     ↪ and validation sets
55     for indices in class_indices.values():
56         np.random.shuffle(indices)
57         split = int(len(indices) * val_percentage)
58         val_indices.extend(indices[:split])
59         train_indices.extend(indices[split:])
60
61     return Subset(dataset, train_indices),
     ↪ Subset(dataset, val_indices)
62
63 # Load the train set
64 trainset = datasets.CIFAR10(root='./data',
     ↪ train=True, download=True, transform=transform)

```

Block	Layer	Type	Parameters
Convolution Block	0	Conv2d	in=3, out=192, kernel=3x3, stride=1, padding=1
	1	BatchNorm2d	features=192, eps=0.001, momentum=0.1
	2	ReLU	
Inception Module	1x1 branch	1x1 Convolution	in=192, out=64, kernel=3x3, stride=1, padding=1
		3x3 Convolution	in=64, out=128, kernel=3x3, stride=1
	3x3 branch	1x1 Convolution	in=192, out=16, kernel=1x1, stride=1
		5x5 Convolution	in=16, out=32, kernel=5x5, stride=1
	5x5 branch	3x3 Max Pooling	kernel=3x3, stride=1, padding=1
		1x1 Convolution	in=192, out=32, kernel=3x3, stride=1
Inception Module	Pooling Branch	1x1 Convolution	in=256, out=128, kernel=3x3, stride=1, padding=1
		3x3 Max Pooling	kernel=3x3, stride=1, padding=1
	1x1 branch	1x1 Convolution	in=256, out=128, kernel=1x1, stride=1
		3x3 Convolution	in=128, out=192, kernel=3x3, stride=1
	3x3 branch	1x1 Convolution	in=256, out=32, kernel=1x1, stride=1
		5x5 Convolution	in=32, out=96, kernel=5x5, stride=1
Inception Module	Pooling Branch	3x3 Max Pooling	kernel=3x3, stride=1, padding=1
		1x1 Convolution	in=256, out=64, kernel=3x3, stride=1
	Max Pooling	0	Max Pooling
			kernel=3x3, stride=2, padding=1
	Inception Module	1x1 branch	in=480, out=192, kernel=3x3, stride=1, padding=1
		3x3 branch	in=480, out=96, kernel=1x1, stride=1
		5x5 branch	in=96, out=208, kernel=3x3, stride=1
		Pooling Branch	in=480, out=16, kernel=1x1, stride=1
		1x1 Convolution	in=16, out=48, kernel=5x5, stride=1
		3x3 Max Pooling	kernel=3x3, stride=1, padding=1
Inception Module	1x1 branch	1x1 Convolution	in=480, out=64, kernel=3x3, stride=1
		3x3 Convolution	kernel=3x3, stride=1, padding=1
	3x3 branch	1x1 Convolution	in=512, out=160, kernel=3x3, stride=1, padding=1
		5x5 Convolution	in=160, out=112, kernel=1x1, stride=1
	5x5 branch	1x1 Convolution	in=512, out=224, kernel=3x3, stride=1
		5x5 Convolution	in=224, out=24, kernel=1x1, stride=1
Inception Module	Pooling Branch	3x3 Max Pooling	in=512, out=64, kernel=5x5, stride=1
		1x1 Convolution	kernel=3x3, stride=1, padding=1
	1x1 branch	1x1 Convolution	in=512, out=192, kernel=3x3, stride=1, padding=1
		3x3 Convolution	in=192, out=128, kernel=1x1, stride=1
	3x3 branch	1x1 Convolution	in=512, out=256, kernel=3x3, stride=1
		5x5 Convolution	in=128, out=24, kernel=1x1, stride=1
Inception Module	Pooling Branch	3x3 Max Pooling	in=512, out=64, kernel=5x5, stride=1
		1x1 Convolution	kernel=3x3, stride=1, padding=1
	1x1 branch	1x1 Convolution	in=512, out=112, kernel=3x3, stride=1, padding=1
		3x3 Convolution	in=112, out=144, kernel=1x1, stride=1
	3x3 branch	1x1 Convolution	in=512, out=288, kernel=3x3, stride=1
		5x5 Convolution	in=144, out=32, kernel=1x1, stride=1
Inception Module	Pooling Branch	3x3 Max Pooling	in=512, out=64, kernel=5x5, stride=1
		1x1 Convolution	kernel=3x3, stride=1, padding=1
	1x1 branch	1x1 Convolution	in=512, out=112, kernel=3x3, stride=1, padding=1
		3x3 Convolution	in=112, out=144, kernel=1x1, stride=1
	3x3 branch	1x1 Convolution	in=512, out=288, kernel=3x3, stride=1
		5x5 Convolution	in=144, out=32, kernel=1x1, stride=1
Inception Module	Pooling Branch	3x3 Max Pooling	in=512, out=64, kernel=5x5, stride=1
		1x1 Convolution	kernel=3x3, stride=1, padding=1
	1x1 branch	1x1 Convolution	in=528, out=256, kernel=3x3, stride=1, padding=1
		3x3 Convolution	in=256, out=160, kernel=1x1, stride=1
	3x3 branch	1x1 Convolution	in=528, out=320, kernel=3x3, stride=1
		5x5 Convolution	in=160, out=32, kernel=1x1, stride=1
	5x5 branch	1x1 Convolution	in=528, out=128, kernel=5x5, stride=1
		5x5 Convolution	in=32, out=128, kernel=3x3, stride=1
Inception Module	Pooling Branch	3x3 Max Pooling	kernel=3x3, stride=1, padding=1
		1x1 Convolution	in=528, out=128, kernel=3x3, stride=1

Table 8: Inception Architecture

Max Pooling	0	Max Pooling	kernel=3x3, stride=2, padding=1
Inception Module	1x1 branch	1x1 Convolution	in=832, out=256, kernel=3x3, stride=1, padding=1
		3x3 Convolution	in=832, out=160, kernel=1x1, stride=1
	3x3 branch	1x1 Convolution	in=160, out=320, kernel=3x3, stride=1
		3x3 Convolution	in=832, out=32, kernel=1x1, stride=1
	5x5 branch	1x1 Convolution	in=32, out=128, kernel=5x5, stride=1
		5x5 Convolution	in=832, out=128, kernel=3x3, stride=1
	Pooling Branch	3x3 Max Pooling	kernel=3x3, stride=1, padding=1
		1x1 Convolution	in=832, out=128, kernel=3x3, stride=1
	Adaptive Average Pooling	0	Adaptive Average Pooling
	Dropout	p=0.2	
Fully Connected Block	0	Linear	in=1024, out=10
	1	ReLU	

Table 9: Inception Architecture Continued

```

65 # Split the dataset ensuring each class is equally 101
66   ↪ represented
67 train_data, val_data = split_by_class(trainset, 102
68   ↪ val_percentage=0.1) 103
69
70 # Load the test set 104
71 testset = datasets.CIFAR10(root='./data', 105
72   ↪ train=False, download=True, transform=transform)
73
74 # Load and transform the CIFAR10 dataset for 106
75   ↪ training, validation, and testing. 107
76 trainloader = DataLoader(train_data, batch_size=64, 108
77   ↪ shuffle=True)
78 valloader = DataLoader(val_data, batch_size=64, 109
79   ↪ shuffle=True)
80 testloader = DataLoader(testset, batch_size=64, 110
81   ↪ shuffle=False)
82
83 # Print the number of images in each part of the 111
84   ↪ dataset to verify the splits. 112
85
86 print(len(train_data)) 113
87 print(len(val_data)) 114
88 print(len(testset)) 115
89
90 # Define the VGG16 neural network architecture with 116
91   ↪ convolutional and fully connected layers. 117
92 class VGG16(nn.Module): 118
93     def __init__(self): 119
94       super(VGG16, self).__init__() 120
95       # Define convolutional blocks with batch 121
96       ↪ normalization and max pooling 122
97
98       self.conv_block1 = nn.Sequential( 123
99         nn.Conv2d(in_channels=3, 124
100        ↪ out_channels=64, kernel_size=3, 125
101        ↪ padding=1), 126
102        nn.BatchNorm2d(num_features=64), 127
103        nn.ReLU(), 128
104        nn.MaxPool2d(kernel_size=2, stride=2)) 129
105
106       self.conv_block2 = nn.Sequential( 130
107         nn.Conv2d(in_channels=64, 131
108        ↪ out_channels=128, kernel_size=3, 132
109        ↪ padding=1), 133
110        nn.BatchNorm2d(num_features=128), 134
111        nn.ReLU(), 135
112        nn.MaxPool2d(kernel_size=2, stride=2)) 136
113
114       self.conv_block3 = self._create_block(128, 137
115        ↪ 256) 138
116
117     self.conv_block4 = self._create_block(256, 139
118        ↪ 256, pool=True) 140
119
120     self.conv_block5 = self._create_block(256, 141
121        ↪ 512) 142
122     self.conv_block6 = self._create_block(512, 143
123        ↪ 512, pool=True) 144
124
125     self.conv_block7 = self._create_block(512, 145
126        ↪ 512) 146
127     self.conv_block8 = self._create_block(512, 147
128        ↪ 512, pool=True) 148
129
130     self.fc_block = nn.Sequential( 149
131       nn.Dropout(0), 150
132       nn.Linear(7*7*512, 512), 153
133       nn.ReLU()) 154
134     self.fc_block1 = nn.Sequential( 155
135       nn.Dropout(0), 156
136       nn.Linear(512, 256), 157
137       nn.ReLU()) 158
138     self.fc_block2 = nn.Linear(256, 10) 159
139
140   def forward(self, x): 160
141     # Apply each convolutional block in sequence 161
142     x = self.conv_block1(x) 162
143     x = self.conv_block2(x) 163
144     x = self.conv_block3(x) 164
145     x = self.conv_block4(x) 165
146     x = self.conv_block5(x) 166
147     x = self.conv_block6(x) 167
148     x = self.conv_block7(x) 168
149     x = self.conv_block8(x) 169
150
151     # Flatten the output for the fully 170
152     ↪ connected layer 171
153     x = x.view(x.size(0), -1) 172
154
155     # Apply the fully connected layers 173
156     x = self.fc_block(x) 174
157     x = self.fc_block1(x) 175
158     x = self.fc_block2(x) 176
159     return x 177
160
161   def _create_block(self, in_channels, 178
162     ↪ out_channels, pool=False): 179
163     layers = [ 180
164       nn.Conv2d(in_channels, out_channels, 181
165         ↪ kernel_size=3, padding=1), 182
166       nn.BatchNorm2d(out_channels), 183
167       nn.ReLU()] 184

```

```

144     if pool:
145         layers.append(nn.MaxPool2d(kernel_size=2,
146                                     stride=2))
146     return nn.Sequential(*layers)
147
148 # Instantiate the VGG16 model
149 vgg16_model = VGG16()
150
151 # Print the model architecture
152 print(vgg16_model)
153 # Install and import torchsummary for a detailed
154 # summary of the model.
155
156 from torchsummary import summary
157 summary(vgg16_model, input_size=(3, 32, 32))
158
159 import numpy as np
160 import torch
161 from torch.utils.data import DataLoader,
162     random_split
163 import torch.optim as optim
164 from torchvision import datasets, transforms
165 from tqdm import tqdm
166
167 device = torch.device("cuda:0" if
168     torch.cuda.is_available() else "cpu")
169 print("Using device:", device)
170 # Move the model to the GPU and set up the loss
171 # function and optimizer.
172
173 model = VGG16()
174 model.to(device) # Move the model to the GPU
175
176 # Loss and Optimizer
177 criterion = nn.CrossEntropyLoss()
178 optimizer = optim.SGD(model.parameters(), lr=0.05,
179     momentum=0.0)
180 # Define variables for tracking progress and early
181 # stopping.
182
183 epochs = 30
184 train_loss_values = []
185 val_loss_values = []
186 val_accuracy_values = []
187 best_val_loss = float('inf')
188 epochs_no_improve = 0
189 n_epochs_stop = 10 # number of epochs to continue
# after no improvement
190
191 for epoch in range(epochs):
192     # Training phase
193     model.train()
194     running_loss = 0.0
195     correct_predictions = 0
196     total_predictions = 0
197
198     for i, data in enumerate(trainloader, 0):
199         inputs, labels = data[0].to(device),
# Move data to GPU
200         data[1].to(device)
201         optimizer.zero_grad()
202
203         outputs = model(inputs)
204         _, predicted = torch.max(outputs.data, 1)
205         total_predictions += labels.size(0)
206         correct_predictions += (predicted ==
# labels).sum().item()
207
208         loss = criterion(outputs, labels)
209         loss.backward()
210         optimizer.step()
211
212         running_loss += loss.item()
213
214 epoch_loss = running_loss / len(trainloader)
215 epoch_accuracy = 100 * correct_predictions /
# total_predictions
216 train_loss_values.append(epoch_loss)
217 train_accuracy_values.append(epoch_accuracy)
218
219 # Validation phase
220 model.eval()
221 val_running_loss = 0.0
222 val_correct_predictions = 0
223 val_total_predictions = 0
224
225 with torch.no_grad():
226     for data in valloader:
227         inputs, labels = data[0].to(device),
# data[1].to(device)
228
229         outputs = model(inputs)
230         _, predicted = torch.max(outputs.data,
# 1)
231         val_total_predictions += labels.size(0)
232         val_correct_predictions += (predicted ==
# labels).sum().item()
233
234         loss = criterion(outputs, labels)
235         val_running_loss += loss.item()
236
237 val_loss = val_running_loss / len(valloader)
238 val_accuracy = 100 * val_correct_predictions /
# val_total_predictions
239 val_loss_values.append(val_loss)
240 val_accuracy_values.append(val_accuracy)
241 print(f'Epoch: {epoch + 1}, Training Loss:
# {epoch_loss:.3f}, Training Accuracy:
# {epoch_accuracy:.2f}%, Validation Loss:
# {val_loss:.3f}, Validation Accuracy:
# {val_accuracy:.2f}%)')
242
243 # Plotting training and validation metrics
244
245 # Early stopping
246 if val_loss < best_val_loss:
247     best_val_loss = val_loss
248     epochs_no_improve = 0
249 else:
250     epochs_no_improve += 1
251     if epochs_no_improve == n_epochs_stop:
252         print('Early stopping!')
253         break
254
255 print('Finished Training')
256
257 print(train_accuracy_values)
258
259 plt.subplot(1, 2, 1)
260 plt.plot(range(1, epoch+2), train_loss_values,
# marker='o', linestyle='-', color='b')
261 plt.xlabel('Epoch')
262 plt.ylabel('Training Loss')
263 plt.title('Training Loss vs. Epoch')
264
265 plt.subplot(1, 2, 2)
266 plt.plot(range(1, epoch + 2),
# train_accuracy_values, marker='o',
# linestyle='-', color='r')
267 plt.xlabel('Epoch')
268 plt.ylabel('Train Accuracy')
269 plt.title('Train Accuracy vs. Epoch')
270
271 plt.tight_layout()
272 plt.show()
273
274 from sklearn.metrics import confusion_matrix
275 import matplotlib.pyplot as plt
276 import seaborn as sns

```

```

275 correct = 0
276 total = 0
277 # Set the model to evaluation mode
278 model.eval()
279
280 # Initialize lists to store true and predicted
281 # labels
282 true_labels = []
283 predicted_labels = []
284
285 with torch.no_grad():
286     for i, data in enumerate(testloader, 0):
287         inputs, labels = data[0].to(device),
288         # data[1].to(device) # Move data to GPU
289
290         # Forward pass
291         outputs = model(inputs)
292
293         # Get predicted labels
294         _, predicted = torch.max(outputs.data, 1)
295
296         # Append true and predicted labels to the
297         # lists
298         true_labels.extend(labels.cpu().numpy())
299         predicted_labels.extend(predicted.cpu().numpy())
300         total += labels.size(0)
301         correct += (predicted ==
302             # labels).sum().item()
303
304 test_accuracy = correct / total
305 print(f'Test Accuracy: {test_accuracy * 100:.2f}%')
306
307 from sklearn.metrics import confusion_matrix
308 import matplotlib.pyplot as plt
309 import seaborn as sns
310 import numpy as np
311
312 def calculate_confusion_matrix(dataloader):
313     model.eval() # Set the model to evaluation mode
314     all_preds = []
315     all_labels = []
316     with torch.no_grad():
317         for data in dataloader:
318             inputs, labels = data[0].to(device),
319             # data[1].to(device)
320             outputs = model(inputs)
321             _, predicted = torch.max(outputs.data,
322             # 1)
323             all_preds.extend(predicted.cpu().numpy())
324             all_labels.extend(labels.cpu().numpy())
325
326         return confusion_matrix(all_labels, all_preds)
327
328 # Calculate confusion matrices for the datasets
329 train_confusion_matrix =
330     # calculate_confusion_matrix(trainloader)
331 val_confusion_matrix =
332     # calculate_confusion_matrix(valloader)
333 test_confusion_matrix =
334     # calculate_confusion_matrix(testloader)
335
336 def plot_all(train_cm, val_cm, test_cm,
337     # train_loss_values, train_accuracy_values,
338     # val_accuracy_values, val_loss_values, epochs):
339     class_names = ['Airplane', 'Automobile',
340     # 'Bird', 'Cat', 'Deer', 'Dog', 'Frog',
341     # 'Horse', 'Ship', 'Truck']
342
343     fig, axes = plt.subplots(2, 3, figsize=(30,
344     # 20)) # 2 rows, 3 columns
345     # Define plots for confusion matrices and
346     # performance metrics
347
348     # Plotting confusion matrices
349
350     sns.heatmap(train_cm, annot=True, fmt="d",
351     # ax=axes[0, 0], xticklabels=class_names,
352     # yticklabels=class_names)
353     axes[0, 0].set_title('Training Set Confusion
354     Matrix')
355     axes[0, 0].set_ylabel('Actual Label')
356     axes[0, 0].set_xlabel('Predicted Label')
357
358     sns.heatmap(val_cm, annot=True, fmt="d",
359     # ax=axes[0, 1], xticklabels=class_names,
360     # yticklabels=class_names)
361     axes[0, 1].set_title('Validation Set Confusion
362     Matrix')
363     axes[0, 1].set_ylabel('Actual Label')
364     axes[0, 1].set_xlabel('Predicted Label')
365
366     sns.heatmap(test_cm, annot=True, fmt="d",
367     # ax=axes[0, 2], xticklabels=class_names,
368     # yticklabels=class_names)
369     axes[0, 2].set_title('Test Set Confusion
370     Matrix')
371     axes[0, 2].set_ylabel('Actual Label')
372     axes[0, 2].set_xlabel('Predicted Label')
373
374     # Plotting training loss
375     axes[1, 0].plot(range(1, epochs + 2),
376     # train_loss_values, marker='o',
377     # linestyle='-', color='b', label='Training
378     # Loss')
379     axes[1, 0].set_xlabel('Epoch')
380     axes[1, 0].set_ylabel('Loss')
381     axes[1, 0].set_title('Training Loss vs. Epoch')
382     axes[1, 0].legend()
383
384     # Plotting training and validation accuracy
385     axes[1, 1].plot(range(1, epochs + 2),
386     # train_accuracy_values, marker='o',
387     # linestyle='-', color='r', label='Training
388     # Accuracy')
389     axes[1, 1].plot(range(1, epochs + 2),
390     # val_accuracy_values, marker='x',
391     # linestyle='-', color='g', label='Validation
392     # Accuracy')
393     axes[1, 1].set_xlabel('Epoch')
394     axes[1, 1].set_ylabel('Accuracy')
395     axes[1, 1].set_title('Accuracy vs. Epoch')
396     axes[1, 1].legend()
397
398     # Plotting validation loss
399     axes[1, 2].plot(range(1, epochs + 2),
400     # val_loss_values, marker='s', linestyle='--',
401     # color='m', label='Validation Loss')
402     axes[1, 2].set_xlabel('Epoch')
403     axes[1, 2].set_ylabel('Loss')
404     axes[1, 2].set_title('Validation Loss vs.
405     Epoch')
406     axes[1, 2].legend()
407
408     plt.tight_layout()
409     plt.show()
410
411
412 plot_all(train_confusion_matrix,
413     # val_confusion_matrix, test_confusion_matrix,
414     # train_loss_values, train_accuracy_values,
415     # val_accuracy_values, val_loss_values, epoch)

```

C. ResNet Python Implementation

```

1 import torch
2 import torch.nn as nn
3 import torch.optim as optim
4 import torchvision.transforms as transforms

```

```

6   from torch.utils.data import DataLoader,
    ↵ random_split
7   from torchvision.datasets import CIFAR10
8
9
10  class BasicBlock(nn.Module):
11      """
12          BasicBlock: A basic building block for the
13          ↵ ResNet architecture.
14
15          Args:
16              in_channels (int): Number of input channels
17                  ↵
18              out_channels (int): Number of output
19                  ↵ channels.
20              stride (int, optional): The stride used for
21                  ↵ the first convolutional layer (default:
22                  ↵ 1).
23
24          Attributes:
25              - conv1: First convolutional layer.
26              - bnl: Batch normalization after the first
27                  ↵ convolution.
28              - relu: ReLU activation function.
29              - conv2: Second convolutional layer.
30              - bn2: Batch normalization after the second
31                  ↵ convolution.
32              - shortcut: Shortcut connection to match
33                  ↵ dimensions if needed.
34
35          Methods:
36              - forward(x): Forward pass through the
37                  ↵ basic block.
38
39          Notes:
40              - The residual connection is added to the
41                  ↵ output before the final ReLU activation
42              - If the stride is not 1 or the number of
43                  ↵ input channels doesn't match the output
44                  ↵ channels,
45                  ↵ a shortcut connection is added to match
46                  ↵ dimensions.
47
48          References:
49              - Original ResNet Paper:
50                  ↵ https://arxiv.org/abs/1512.03385
51              - torchvision.models.resnet implementation:
52                  ↵ https://pytorch.org/vision/stable/\_modu
53                  ↵ les/torchvision/models/resnet.html
54
55      """
56
57      expansion = 1
58
59      def __init__(self, in_channels, out_channels,
60          ↵ stride=1):
61          super(BasicBlock, self).__init__()
62          self.conv1 = nn.Conv2d(in_channels,
63              ↵ out_channels, kernel_size=3,
64              ↵ stride=stride, padding=1, bias=False)
65          self.bn1 = nn.BatchNorm2d(out_channels)
66          self.relu = nn.ReLU(inplace=True)
67          self.conv2 = nn.Conv2d(out_channels,
68              ↵ out_channels, kernel_size=3, stride=1,
69              ↵ padding=1, bias=False)
70          self.bn2 = nn.BatchNorm2d(out_channels)
71
72          self.shortcut = nn.Sequential()
73          if stride != 1 or in_channels !=
74              ↵ out_channels:
75              self.shortcut = nn.Sequential(
76                  nn.Conv2d(in_channels,
77                      ↵ out_channels, kernel_size=1,
78                      ↵ stride=stride, bias=False),
79                  nn.BatchNorm2d(out_channels))
80
81      def forward(self, x):
82
83          residual = x
84          out = self.conv1(x)
85          out = self.bn1(out)
86          out = self.relu(out)
87          out = self.conv2(out)
88          out = self.bn2(out)
89          out += self.shortcut(residual)
90          out = self.relu(out)
91
92      return out
93
94
95
96      class ResNet(nn.Module):
97          """
98              ResNet: Implementation of the ResNet
99                  ↵ architecture.
100
101              Args:
102                  block (nn.Module): The basic building block
103                      ↵ for the ResNet architecture.
104                  layers (list): List of integers
105                      ↵ representing the number of blocks in
106                      ↵ each layer. The length of the list
107                          determines the total number
108                          ↵ of layers in the network.
109                  num_classes (int, optional): Number of
110                      ↵ classes in the classification task
111                      ↵ (default: 10).
112                  dropout_prob (float, optional): Probability
113                      ↵ of dropout for regularization (default:
114                      ↵ 0.5).
115
116              Attributes:
117                  - conv1: Initial convolutional layer.
118                  - bn1: Initial batch normalization layer.
119                  - relu: ReLU activation function.
120                  - layer1, layer2, layer3, layer4: Residual
121                      ↵ layers.
122                  - avg_pool: Adaptive average pooling layer.
123                  - fc: Fully connected layer for
124                      ↵ classification.
125                  - dropout: Dropout layer for regularization.
126
127              Methods:
128                  - forward(x): Forward pass through the
129                      ↵ network.
130
131              Notes:
132                  - This implementation follows the ResNet
133                      ↵ architecture with configurable layers
134                      ↵ and block types.
135                  - The network structure is based on the
136                      ↵ ResNet18 and ResNet34 variants with
137                      ↵ different numbers of layers.
138                  - For ResNet18, set layers=[2, 2, 2, 2].
139                  - For ResNet34, set layers=[3, 4, 6, 3].
140
141              References:
142                  - Original ResNet Paper:
143                      ↵ https://arxiv.org/abs/1512.03385
144                  - torchvision.models.resnet implementation:
145                      ↵ https://pytorch.org/vision/stable/\_modu
146                      ↵ les/torchvision/models/resnet.html
147
148      """
149
150      def __init__(self, block, layers,
151          ↵ num_classes=10, dropout_prob=0.5):
152          super(ResNet, self).__init__()
153          self.in_channels = 64
154          self.conv1 = nn.Conv2d(3, 64,
155              ↵ kernel_size=3, stride=1, padding=1,
156              ↵ bias=False)
157          self.bn1 = nn.BatchNorm2d(64)
158          self.relu = nn.ReLU(inplace=True)
159          self.layer1 = self._make_layer(block, 64,
160              ↵ layers[0], stride=1)
161          self.layer2 = self._make_layer(block, 128,
162              ↵ layers[1], stride=2)

```

```

110     self.layer3 = self._make_layer(block, 256, 166
111         ↪ layers[2], stride=2) 167
112     self.layer4 = self._make_layer(block, 512, 168
113         ↪ layers[3], stride=2) 169
114     self.avg_pool = nn.AdaptiveAvgPool2d((1,1))170
115     self.fc = nn.Linear(512 * block.expansion, 171
116         ↪ num_classes) 172
117     self.dropout = nn.Dropout(p=dropout_prob) 172
118
119     def _make_layer(self, block, out_channels, 174
120         ↪ blocks, stride): 174
121         layers = [block(self.in_channels, 175
122             ↪ out_channels, stride)] 175
123         self.in_channels = out_channels * 176
124             ↪ block.expansion 176
125         for _ in range(1, blocks): 177
126             layers.append(block(self.in_channels, 177
127                 ↪ out_channels, stride=1)) 178
128         return nn.Sequential(*layers) 179
129
130     def forward(self, x):
131         x = self.conv1(x) 180
132         x = self.bn1(x) 181
133         x = self.relu(x) 182
134         x = self.layer1(x) 182
135         x = self.layer2(x) 183
136         x = self.layer3(x) 183
137         x = self.layer4(x) 184
138         x = self.avg_pool(x) 184
139         x = x.view(x.size(0), -1) 185
140         x = self.dropout(x) 185
141         x = self.fc(x) 186
142         return x 186
143
144     def test_model(model, test_loader, device):
145         """
146             Evaluate the performance of a trained model on 188
147             ↪ a test dataset. 189
148
149             Args:
150                 model (nn.Module): The trained model to be 190
151                     ↪ evaluated. 191
152                 test_loader (DataLoader): DataLoader for 192
153                     ↪ the test dataset. 193
154                 device (torch.device): The device (cuda or 194
155                     ↪ cpu) on which the evaluation should be 195
156                     ↪ performed. 196
157
158             Notes:
159                 - The model is set to evaluation mode using 198
160                     ↪ `model.eval()`. 199
161                 - The function iterates through the test 200
162                     ↪ dataset, computes predictions, and
163                     ↪ calculates accuracy.
164                 - The accuracy is printed to the console.
165                 - Ensure that the model and test_loader are 201
166                     ↪ properly configured and loaded before
167                     ↪ calling this function.
168                 - The device argument determines whether 202
169                     ↪ the evaluation is performed on GPU or
170                     ↪ CPU.
171
172             """
173             model.eval()
174             correct = 0 204
175             total = 0 205
176
177             with torch.no_grad():
178                 for images, labels in test_loader:
179                     images, labels = images.to(device),
180                         ↪ labels.to(device) 209
181                     outputs = model(images)
182                     _, predicted = torch.max(outputs, 1) 210
183                     total += labels.size(0) 211
184                     correct += (predicted ==
185                         ↪ labels).sum().item() 212
186
187             accuracy = correct / total
188             print(f'Test Accuracy: {accuracy * 100:.2f}%')
189
190     def main():
191         """
192             Main function for training and evaluating a
193             ↪ ResNet model on the CIFAR-10 dataset.
194
195             This function performs the following steps:
196             1. Initializes a ResNet model (ResNet18) with
197                 ↪ specified configurations.
198             2. Sets up data transformations and creates
199                 ↪ training and validation datasets.
200             3. Defines the loss function (CrossEntropyLoss)
201                 ↪ and optimizer (SGD) for training.
202             4. Iterates through epochs, training the model
203                 ↪ and evaluating on the validation set.
204             5. Prints training and validation metrics
205                 ↪ during each epoch.
206             6. Evaluates the trained model on the test set
207                 ↪ and generates confusion matrices.
208             7. Plots confusion matrices and
209                 ↪ training/validation curves.
210
211             Notes:
212                 - Ensure that the required modules (torch,
213                     ↪ torchvision, etc.) are installed.
214                 - The ResNet model, dataset paths, and
215                     ↪ hyperparameters can be modified as
216                     ↪ needed.
217                 - Confusion matrices and training curves
218                     ↪ are visualized using the plot_all
219                     ↪ function.
220                 - The model is trained on GPU if available,
221                     ↪ otherwise, it falls back to CPU.
222
223             """
224
225             device = torch.device("cuda" if
226                 ↪ torch.cuda.is_available() else "cpu")
227             resnet18 = ResNet(BasicBlock, [2, 2, 2, 2],
228                 ↪ dropout_prob=0.5).to(device)
229
230
231             batch_size = 64
232             learning_rate = 0.001
233             momentum = 0
234             epochs = 25
235
236
237             transform = transforms.Compose([transforms.RandJ
238                 ↪ omHorizontalFlip(),
239                     ↪ transforms.RandJ
240                         ↪ omCrop(32,
241                             ↪ padding=4),
242                             ↪ transforms.ToTeJ
243                                 ↪ nsor(),
244                                 ↪ transforms.NormJ
245                                     ↪ alize((0.5,
246                                         ↪ 0.5, 0.5),
247                                             ↪ (0.5, 0.5,
248                                                 ↪ 0.5))])
249
250
251             train_dataset = CIFAR10(root='./data',
252                 ↪ train=True, download=True,
253                     ↪ transform=transform)
254             torch.manual_seed(0)
255
256
257             trainset_size = len(train_dataset)
258             valset_size = int(0.1 * trainset_size)
259             trainset_size = trainset_size - valset_size
260
261
262             train_subset, val_subset =
263                 ↪ random_split(train_dataset, [trainset_size,
264                     ↪ valset_size])
265
266
267             train_loader = DataLoader(train_subset,
268                 ↪ batch_size=batch_size, shuffle=True,
269                     ↪ num_workers=2)

```

```

213     valloader = DataLoader(val_subset,
214         batch_size=64, shuffle=False)
215
216     criterion = nn.CrossEntropyLoss()
217     optimizer = optim.SGD(resnet18.parameters(),
218         lr=learning_rate, momentum=momentum)
219
220     train_loss_values = []
221     train_accuracy_values = []
222     val_loss_values = []
223     val_accuracy_values = []
224
225     for epoch in range(epochs):
226
226         resnet18.train()
227         total_train_correct = 0
228         total_train_samples = 0
229         total_train_loss = 0.0
230
231         for i, (images, labels) in
232             enumerate(train_loader):
233             images, labels = images.to(device),
234                 labels.to(device)
235             optimizer.zero_grad()
236             outputs = resnet18(images)
237             loss = criterion(outputs, labels)
238             loss.backward()
239             optimizer.step()
240
241             _, predicted = torch.max(outputs, 1)
242             total_train_samples += labels.size(0)
243             total_train_correct += (predicted ==
244                 labels).sum().item()
245             total_train_loss += loss.item()
246
247             if (i + 1) % 100 == 0:
248                 print(f'Training - Epoch
249                     [{epoch+1}/{epochs}], Step
250                     [{i+1}/{len(train_loader)}],
251                     Loss: {loss.item():.4f}')
252
253     train_accuracy = total_train_correct /
254         total_train_samples
255     train_loss = total_train_loss /
256         len(train_loader)
257     train_conf_matrix =
258         calculate_confusion_matrix(resnet18,
259             train_loader, device)
260
261     resnet18.eval()
262     total_val_correct = 0
263     total_val_samples = 0
264     total_val_loss = 0.0
265
266     with torch.no_grad():
267         for images, labels in valloader:
268             images, labels = images.to(device),
269                 labels.to(device)
270             outputs = resnet18(images)
271             loss = criterion(outputs, labels)
272             _, predicted = torch.max(outputs, 1)
273             total_val_samples += labels.size(0)
274             total_val_correct += (predicted ==
275                 labels).sum().item()
276             total_val_loss += loss.item()
277
278     val_accuracy = total_val_correct /
279         total_val_samples
280     val_loss = total_val_loss / len(valloader)
281     val_conf_matrix =
282         calculate_confusion_matrix(resnet18,
283             valloader, device)
284
285     print(f'Epoch [{epoch+1}/{epochs}] -
286         Training Accuracy: {train_accuracy * 100:.2f}%, Validation Accuracy:
287             {val_accuracy * 100:.2f}%)')
288     train_loss_values.append(train_loss)
289     train_accuracy_values.append(train_accuracy)
290     val_loss_values.append(val_loss)
291     val_accuracy_values.append(val_accuracy)
292
293     print("Training finished.")
294
295     test_batch_size = 10
296     test_dataset = CIFAR10(root='./data',
297         train=False, download=True,
298         transform=transform)
299     test_loader = DataLoader(test_dataset,
300         batch_size=test_batch_size, shuffle=False,
301         num_workers=2)
302
303     resnet18.eval()
304     test_model(resnet18, test_loader, device)
305     test_conf_matrix =
306         calculate_confusion_matrix(resnet18,
307             test_loader, device)
308     plot_all(train_conf_matrix, val_conf_matrix,
309             test_conf_matrix, train_loss_values,
310             train_accuracy_values,
311             val_accuracy_values, val_loss_values,
312             epochs, model_no=18, show=True)
313
314     if __name__ == '__main__':
315         main()

```

D. DenseNet Python Implementation

```

1 import numpy as np
2 import torch
3 import torch.nn as nn
4 from torch.utils.data import DataLoader,
5     random_split
6 import torch.optim as optim
7 from torchvision import datasets, transforms
8 from tqdm import tqdm
9 import matplotlib.pyplot as plt
10 import seaborn as sns
11 import numpy as np
12 import copy
13 from sklearn.metrics import confusion_matrix
14 from tqdm import tqdm
15 import csv
16 from torchsummary import summary
17 from collections import OrderedDict
18
19 # Transition Layer to connect Dense Layers
20 class Transition(nn.Sequential):
21     """
22     Transition layer used in DenseNet between dense
23     blocks.
24
25     This layer performs convolution and pooling
26     operations to reduce the size of the
27     feature maps.
28
29     Parameters:
30         in_channels (int): Number of input channels.
31         compression_factor (float): Compression factor
32             to reduce the feature map size.
33         version (str): Specifies the version of the
34             transition layer. 'Default' version
35                 keeps the number of channels
36                 unchanged, while 'C' applies
37                 compression.
38     """
39
40

```

```

32     def __init__(self, in_channels,           80
33         compression_factor=0, version='Default'):
34         super(Transition, self).__init__()
35         if version == 'Default':
36             self.add_module('batchnorm',          81
37                 nn.BatchNorm2d(in_channels))  82
38             self.add_module('conv1x1',            83
39                 nn.Conv2d(in_channels, in_channels, 84
40                     kernel_size=1, bias=False)) 85
41             self.add_module('avgpool2x2',        86
42                 nn.AvgPool2d(kernel_size=2))   87
43
44         if version == 'C':
45             output_channel = int(in_channels * 88
46                 compression_factor // 1)      89
47
48             self.add_module('batchnorm',          90
49                 nn.BatchNorm2d(in_channels))  91
50             self.add_module('conv1x1',            92
51                 nn.Conv2d(in_channels,          93
52                     output_channel, kernel_size=1, 94
53                     bias=False))           95
54             self.add_module('avgpool2x2',        96
55                 nn.AvgPool2d(kernel_size=2))   97
56
57 # The Dense Block is the main contribution of the
58 # paper
59 class DenseBlock(nn.Module):
60     """
61     DenseBlock in DenseNet, consisting of batch
62     normalization, ReLU activation, and
63     convolution.
64
65     This block forms the core building block of
66     DenseNet, with each DenseBlock layer
67     receiving additional feature maps from all
68     preceding layers.
69
70     Parameters:
71     in_channels (int): Number of input channels.
72     growth_rate (int): How many filters to add per
73     dense layer (growth rate).
74     version (str): Version of the DenseBlock.
75     → 'Default' is the standard block,
76     → while 'B' includes bottleneck
77     → layers.
78
79     def __init__(self, in_channels, growth_rate,
80         version='Default'):
81         super(DenseBlock, self).__init__()
82
83         if version == 'Default':
84             # Default Dense Block
85             self.dense_block =           117
86                 nn.Sequential(OrderedDict([
87                     ('batchnorm0',              118
88                         nn.BatchNorm2d(in_channels)), 119
89                     ('relu0', nn.ReLU(inplace=True)), 120
90                     ('conv3x3', nn.Conv2d(in_channels, 121
91                         growth_rate, kernel_size=3, 122
92                         padding=1, bias=False))    123
93                 ]))                           124
94
95         if version == 'B':
96             # B type Dense Block
97             # Bottleneck layers
98             self.dense_block =           128
99                 nn.Sequential(OrderedDict([
100                     ('batchnorm0',              130
101                         nn.BatchNorm2d(in_channels)), 131
102                     ('relu0', nn.ReLU(inplace=True)), 132
103                     ('conv1x1', nn.Conv2d(in_channels, 134
104                         4 * growth_rate, kernel_size=1, 135
105                         bias=False)),           136
106                     ('batchnorm1', nn.BatchNorm2d(4 * 137
107                         growth_rate)),       138
108                     ('relu1', nn.ReLU(inplace=True)), 139
109                 ])))
110
111     def forward(self, x):
112         """
113         Forward pass of the DenseBlock.
114
115         Concatenates the input with the output
116         → feature maps of the dense block.
117
118         Parameters:
119         x (Tensor): Input tensor to the DenseBlock.
120
121         Returns:
122         Tensor: Concatenated output after passing
123             through the block.
124
125         # Forwards the convolutional layers, and
126         → concatenate it with the input.
127         return torch.cat([x, self.dense_block(x)], 128
128             1)
129
130 # The class below defines cascaded Dense Blocks
131 class DenseLayer(nn.Module):
132     """
133     Layer of DenseNet that consists of multiple
134     → DenseBlocks.
135
136     Parameters:
137     in_channels (int): Number of input channels.
138     growth_rate (int): Growth rate of the dense
139     → blocks.
140     num_blocks (int): Number of dense blocks in the
141     → layer.
142     version (str): Version of DenseBlocks used in
143     → the layer.
144
145     def __init__(self, in_channels, growth_rate,
146         num_blocks, version='Default'):
147         super(DenseLayer, self).__init__()
148
149         blocks = []
150         for i in range(num_blocks):
151             if version == 'Default':
152                 block = DenseBlock(in_channels + i
153                     * growth_rate, growth_rate,
154                     version)
155             elif version == 'B':
156                 block = DenseBlock(in_channels + i
157                     * growth_rate, growth_rate,
158                     version)
159             blocks.append(block)
160
161         self.blocks = nn.Sequential(*blocks)
162
163     def forward(self, x):
164         """
165         Forward pass of the DenseLayer.
166
167         Parameters:
168         x (Tensor): Input tensor to the DenseLayer.
169
170         Returns:
171         Tensor: Output tensor after passing through
172             all dense blocks.
173
174         for block in self.blocks:
175             x = block(x)
176
177         return x
178
179 # The main network connecting the above methods.
180 class DenseNet(nn.Module):
181     """
182
183     DenseNet architecture for image classification.

```

```

140     Parameters:                                     194
141     growth_rate (int): Number of filters to add per 195
142         ↪ dense layer (growth rate).                196
143     compression_factor (float): Compression factor 197
144         ↪ for Transition layers.                  198
145     layer_widths (List[int]): Number of blocks in      199
146         ↪ each dense layer.                      200
147     num_classes (int): Number of output classes for 201
148         ↪ classification.                      202
149     dropout_rate (float): Dropout rate for the       203
150         ↪ final fully connected layer.             204
151     version (str): Version of DenseNet ('Default', 205
152         ↪ 'B', 'BC', etc.).                     206
153     """
154
155     def __init__(self, growth_rate, 207
156         ↪ compression_factor, layer_widths, 208
157         ↪ num_classes, dropout_rate=0.0, 209
158         ↪ version='Default'): 210
159         super(DenseNet, self).__init__() 211
160
161         # Assuming the input is RGB images 212
162         num_channels = 3 213
163
164         dense_version = 'Default' 214
165         transition_version = 'Default' 215
166         comp_fac = 1 216
167
168         if (version == 'B') or (version == 'BC'): 217
169             dense_version = 'B' 218
170
171         if (version == 'BC') or (version == 'C'): 219
172             transition_version = 'C' 220
173             comp_fac = compression_factor 221
174
175         # Initial convolution 222
176         self.features = nn.Sequential(OrderedDict([223
177             ('conv0', nn.Conv2d(num_channels, 224
178                 ↪ growth_rate * 2, kernel_size=3, 225
179                 ↪ stride=1, padding=1, bias=False)), 226
180             ('norm0', nn.BatchNorm2d(growth_rate * 227
181                 ↪ 2)), 228
182             ('relu0', nn.ReLU(inplace=True)) 229
183         ])) 230
184
185         num_channels = growth_rate * 2 231
186
187         # Dense blocks and transition layers 232
188         for i, num_blocks in 233
189             enumerate(layer_widths): 234
190                 # Add DenseBlock 235
191                 dense_block = DenseLayer(num_channels, 236
192                     ↪ growth_rate, num_blocks, 237
193                     ↪ dense_version) 238
194                 self.features.add_module('denseblock%d'% 239
195                     ↪ % (i + 1), dense_block) 240
196
197                 num_channels += num_blocks * growth_rate 241
198
199                 if i != len(layer_widths) - 1: 242
200                     # Add Transition Layer 243
201                     trans_layer = 244
202                         Transition(num_channels, 243
203                             ↪ comp_fac, transition_version) 244
204                         self.features.add_module('transitio 245
205                             ↪ n%d'% (i + 1), 244
206                             ↪ trans_layer) 246
207                         num_channels = int(num_channels * 245
208                             ↪ comp_fac) 247
209
210                     # Final batch norm 248
211                     self.features.add_module('norm5', 249
212                         ↪ nn.BatchNorm2d(num_channels)) 249
213
214                     # Global average pooling 248
215                     self.global_avg_pool = 249
216                         nn.AdaptiveAvgPool2d((1, 1)) 249
217
218
219             # Classifier 240
220             self.classifier = 241
221                 nn.Sequential(OrderedDict([ 242
222                     ('dropout', nn.Dropout(p=dropout_rate)), 243
223                     ('fc', nn.Linear(num_channels, 244
224                         ↪ num_classes)) 245
225                 ])) 246
226
227             def _initialize_weights(self): 247
228                 """
229                     Initialize weights of the network using 248
230                     specific schemes for different types of 249
231                     layers. 249
232
233                     for m in self.modules(): 240
234                         if isinstance(m, nn.Conv2d): 241
235                             nn.init.kaiming_normal_(m.weight, 242
236                             mode='fan_out', 243
237                             nonlinearity='relu') 244
238                         if m.bias is not None: 245
239                             nn.init.constant_(m.bias, 0) 246
240                         elif isinstance(m, nn.BatchNorm2d): 247
241                             nn.init.constant_(m.weight, 1) 248
242                             nn.init.constant_(m.bias, 0) 249
243                         elif isinstance(m, nn.Linear): 240
244                             nn.init.normal_(m.weight, 0, 0.01) 241
245                             nn.init.constant_(m.bias, 0) 246
246
247             def forward(self, x): 247
248                 """
249                     Forward pass of the DenseNet. 250
250
251                     Parameters: 251
252                         x (Tensor): Input tensor to the network. 253
253
254                     Returns: 254
255                         Tensor: Output tensor after classification. 256
256
257                         x = self.features(x) 257
258                         x = self.global_avg_pool(x) 259
259                         x = x.view(x.size(0), -1) 260
260                         x = self.classifier(x) 261
261                         return x 262
262
263
264             def train_model(model, trainloader, valloader, 263
265                 optimizer, criterion, device, epochs, patience, 266
266                 model_no): 267
267                 """
268                     Train a model with given data and 269
269                     hyperparameters, implementing early 270
270                     stopping. 271
271
272                     Parameters: 272
273                         model (nn.Module): The neural network model to 274
274                             be trained. 275
275                         trainloader (DataLoader): DataLoader for the 276
276                             training set. 277
277                         valloader (DataLoader): DataLoader for the 278
278                             validation set. 279
279                         optimizer (Optimizer): Optimizer for the 280
280                             training process. 281
281                         criterion (Loss Function): Loss function used 282
282                             during training. 283
283                         device (str): Device to train the model on 284
284                             ('cpu' or 'cuda'). 285
285                         epochs (int): Total number of epochs for 286
286                             training. 287
287                         patience (int): Patience for early stopping 288
288                             (number of epochs to wait without 289
289                             improvement). 290
290                         model_no (int): Identifier for the model, used 291
291                             for saving metrics. 292
292
293                     Returns: 293

```

```

250      tuple: Tuple containing the trained model, and 311             val_running_loss += loss.item()
251      lists of training and validation losses and 312             val_loss = val_running_loss / len(valloader)
252      accuracies. 313             val_accuracy = 100 *
253      """
254      model.to(device) 314             ← val_correct_predictions /
255
256      train_loss_values = [] 315             ← val_total_predictions
257      train_accuracy_values = [] 316             val_loss_values.append(val_loss)
258      val_loss_values = [] 317             val_accuracy_values.append(val_accuracy)
259      val_accuracy_values = [] 318
260
261      best_val_loss = float('inf') 319
262      best_model_params = 320
263      ← copy.deepcopy(model.state_dict()) # Save 321
264      ← the initial model parameters 322
265      epochs_no_improve = 0 323
266
267      epoch_metrics = [] 324
268      for epoch in range(epochs): 325
269          # Training phase 326
270          model.train() 327
271          running_loss = 0.0 328
272          correct_predictions = 0 329
273          total_predictions = 0 330
274          for i, (inputs, labels) in 331
275              enumerate(trainloader): 332
276              inputs, labels = inputs.to(device), 333
277              ← labels.to(device) # Move data to 334
278              ← GPU 335
279              optimizer.zero_grad() 336
280
281              outputs = model(inputs) 337
282              _, predicted = torch.max(outputs.data, 338
283              ← 1) 339
284              total_predictions += labels.size(0) 340
285              correct_predictions += (predicted == 341
286              ← labels).sum().item() 342
287
288              loss = criterion(outputs, labels) 343
289              loss.backward() 344
290              optimizer.step() 345
291
292              running_loss += loss.item() 346
293
294              if (i + 1) % 100 == 0: 347
295                  print(f'Iteration {i+1}: Loss = 348
296                  ← {running_loss / (i+1):.4f}')
297
298      epoch_loss = running_loss / len(trainloader) 349
299      epoch_accuracy = 100 * correct_predictions 350
300      ← / total_predictions 351
301      train_loss_values.append(epoch_loss) 352
302      train_accuracy_values.append(epoch_accuracy) 353
303
304      # Validation phase 354
305      model.eval() 355
306      val_running_loss = 0.0 356
307      val_correct_predictions = 0 357
308      val_total_predictions = 0 358
309
310      with torch.no_grad(): 359
311          for inputs, labels in valloader: 360
312              inputs, labels = inputs.to(device), 361
313              ← labels.to(device) 362
314
315              outputs = model(inputs) 363
316              _, predicted = 364
317              ← torch.max(outputs.data, 1) 365
318              val_total_predictions += 366
319              ← labels.size(0) 367
320              val_correct_predictions += 368
321              ← (predicted == 369
322              ← labels).sum().item() 370
323
324              loss = criterion(outputs, labels) 371
325
326      val_running_loss += loss.item() 372
327
328      val_loss = val_running_loss / len(valloader) 373
329      val_accuracy = 100 * 374
330      ← val_correct_predictions / 375
331      ← val_total_predictions 376
332      val_loss_values.append(val_loss) 377
333      val_accuracy_values.append(val_accuracy) 378
334
335      print(f'Epoch {epoch+1}/{epochs}, Train 379
336      ← Loss: {epoch_loss:.4f}, Train Acc: 380
337      ← {epoch_accuracy:.2f}%, Val Loss: 381
338      ← {val_loss:.4f}, Val Acc: 382
339      ← {val_accuracy:.2f}%' 383
340      # Save metrics for the current epoch 384
341      epoch_metrics.append([epoch + 1, 385
342      ← epoch_loss, epoch_accuracy, val_loss, 386
343      ← val_accuracy]) 387
344
345      # Early stopping check 388
346      if val_loss < best_val_loss: 389
347          best_val_loss = val_loss 390
348          best_model_params = 391
349          ← copy.deepcopy(model.state_dict()) 392
350          epochs_no_improve = 0 393
351
352      else: 394
353          epochs_no_improve += 1 395
354          if epochs_no_improve >= patience: 396
355              print('Early stopping!')
356              break 397
357
358      model.load_state_dict(best_model_params) 398
359      print('Finished Training') 399
360
361      # Write metrics to a CSV file 400
362      csv_file = f'training_metrics_{model_no}.csv' 401
363      with open(csv_file, 'w', newline='') as file: 402
364          writer = csv.writer(file) 403
365          writer.writerow(['Epoch', 'Train Loss', 404
366          ← 'Train Accuracy', 'Validation Loss', 405
367          ← 'Validation Accuracy']) 406
368          writer.writerows(epoch_metrics) 407
369
370      return model, train_loss_values, 408
371      ← train_accuracy_values, val_loss_values, 409
372      ← val_accuracy_values 410
373
374      def cal_acc(model, testloader, device): 411
375          """
376          Calculate the accuracy of a trained model on a 412
377          ← given dataset. 413
378
379          Parameters: 414
380          model (nn.Module): The trained neural network 415
381          ← model. 416
382          testloader (DataLoader): DataLoader for the 417
383          ← test set. 418
384          device (str): Device on which the model is 419
385          ← loaded ('cpu' or 'cuda'). 420
386
387          Returns: 421
388          tuple: Tuple containing the test accuracy and 422
389          ← the predicted labels. 423
390          """
391          correct = 0 424
392          total = 0 425
393          # Set the model to evaluation mode 426
394          model.to(device) 427
395          model.eval() 428
396
397          # Initialize lists to store true and predicted 429
398          ← labels 430
399          true_labels = [] 431
400          predicted_labels = [] 432

```

```

369     with torch.no_grad():
370         for i, data in enumerate(testloader, 0):
371             inputs, labels = data[0].to(device),
372             ↪ data[1].to(device) # Move data to
373             ↪ GPU
374
375             # Forward pass
376             outputs = model(inputs)
377
378             # Get predicted labels
379             _, predicted = torch.max(outputs.data,
380             ↪ 1)
381
382             # Append true and predicted labels to
383             ↪ the lists
384             true_labels.extend(labels.cpu().numpy())
385             predicted_labels.extend(predicted.cpu())
386             ↪ .numpy())
387             total += labels.size(0)
388             correct += (predicted ==
389             ↪ labels).sum().item()
390             test_accuracy = correct / total
391
392             return test_accuracy, predicted_labels
393
394 def calculate_confusion_matrix(model, dataloader,
395     ↪ device):
396     """
397     Compute the confusion matrix for a model's
398     ↪ predictions on a dataset.
399
400     Parameters:
401     model (nn.Module): The trained neural network
402     ↪ model.
403     dataloader (DataLoader): DataLoader for the
404     ↪ dataset.
405     device (str): Device on which the model is
406     ↪ loaded ('cpu' or 'cuda').
407
408     Returns:
409     ndarray: Confusion matrix of the model's
410     ↪ predictions.
411     """
412     model.eval() # Set the model to evaluation mode
413     all_preds = []
414     all_labels = []
415     with torch.no_grad():
416         for data in dataloader:
417             inputs, labels = data[0].to(device),
418             ↪ data[1].to(device)
419             outputs = model(inputs)
420             _, predicted = torch.max(outputs.data,
421             ↪ 1)
422             all_preds.extend(predicted.cpu().numpy())
423             ↪ )
424             all_labels.extend(labels.cpu().numpy())
425
426     return confusion_matrix(all_labels, all_preds)
427
428 def plot_all(train_cm, val_cm, test_cm,
429     ↪ train_loss_values, train_accuracy_values,
430     ↪ val_accuracy_values, val_loss_values, epochs,
431     ↪ model_no, show=False):
432     """
433     Plot and save various training and evaluation
434     ↪ metrics including confusion matrices and
435     ↪ accuracy/loss curves.
436
437     Parameters:
438     train_cm (ndarray): Confusion matrix for the
439     ↪ training set.
440     val_cm (ndarray): Confusion matrix for the
441     ↪ validation set.
442
443     test_cm (ndarray): Confusion matrix for the
444     ↪ test set.
445     train_loss_values (list): List of training loss
446     ↪ values for each epoch.
447     train_accuracy_values (list): List of training
448     ↪ accuracy values for each epoch.
449     val_accuracy_values (list): List of validation
450     ↪ accuracy values for each epoch.
451     val_loss_values (list): List of validation loss
452     ↪ values for each epoch.
453     epochs (int): Number of epochs.
454     model_no (int): Identifier for the model, used
455     ↪ for saving the plot.
456     show (bool): If True, display the plot.
457     ↪ Defaults to False.
458
459     class_names = ['Airplane', 'Automobile',
460     ↪ 'Bird', 'Cat', 'Deer', 'Dog', 'Frog',
461     ↪ 'Horse', 'Ship', 'Truck']
462
463     fig, axes = plt.subplots(2, 3, figsize=(30,
464     ↪ 20)) # 2 rows, 3 columns
465
466     # Plotting confusion matrices
467     sns.heatmap(train_cm, annot=True, fmt="d",
468     ↪ ax=axes[0, 0], xticklabels=class_names,
469     ↪ yticklabels=class_names)
470     axes[0, 0].set_title('Training Set Confusion
471     ↪ Matrix')
472     axes[0, 0].set_ylabel('Actual Label')
473     axes[0, 0].set_xlabel('Predicted Label')
474
475     sns.heatmap(val_cm, annot=True, fmt="d",
476     ↪ ax=axes[0, 1], xticklabels=class_names,
477     ↪ yticklabels=class_names)
478     axes[0, 1].set_title('Validation Set Confusion
479     ↪ Matrix')
480     axes[0, 1].set_ylabel('Actual Label')
481     axes[0, 1].set_xlabel('Predicted Label')
482
483     sns.heatmap(test_cm, annot=True, fmt="d",
484     ↪ ax=axes[0, 2], xticklabels=class_names,
485     ↪ yticklabels=class_names)
486     axes[0, 2].set_title('Test Set Confusion
487     ↪ Matrix')
488     axes[0, 2].set_ylabel('Actual Label')
489     axes[0, 2].set_xlabel('Predicted Label')
490
491     # Plotting training loss
492     axes[1, 0].plot(range(1, epochs + 1),
493     ↪ train_loss_values, marker='o',
494     ↪ linestyle='-', color='b', label='Training
495     ↪ Loss')
496     axes[1, 0].set_xlabel('Epoch')
497     axes[1, 0].set_ylabel('Loss')
498     axes[1, 0].set_title('Training Loss vs. Epoch')
499     axes[1, 0].legend()
500
501     # Plotting training and validation accuracy
502     axes[1, 1].plot(range(1, epochs + 1),
503     ↪ train_accuracy_values, marker='o',
504     ↪ linestyle='-', color='r', label='Training
505     ↪ Accuracy')
506     axes[1, 1].plot(range(1, epochs + 1),
507     ↪ val_accuracy_values, marker='x',
508     ↪ linestyle='-', color='g', label='Validation
509     ↪ Accuracy')
510     axes[1, 1].set_xlabel('Epoch')
511     axes[1, 1].set_ylabel('Accuracy')
512     axes[1, 1].set_title('Accuracy vs. Epoch')
513     axes[1, 1].legend()
514
515     # Plotting validation loss
516     axes[1, 2].plot(range(1, epochs + 1),
517     ↪ val_loss_values, marker='s', linestyle='--',
518     ↪ color='m', label='Validation Loss')
519     axes[1, 2].set_xlabel('Epoch')

```

```

470     axes[1, 2].set_ylabel('Loss')
471     axes[1, 2].set_title('Validation Loss vs.
472     ↪ Epoch')
473     axes[1, 2].legend()
474
475     plt.tight_layout()
476     if show:
477         plt.show()
478     plt.savefig(f'model_results_{model_no}.png')
479     ↪ Save the figure to a file
480
481 # Check whether the GPU available
482 device = torch.device("cuda:0" if
483     ↪ torch.cuda.is_available() else "cpu")
484 print("Using device:", device)
485
486 # Defining the models.
487 model1 = DenseNet(growth_rate=24,
488     ↪ compression_factor=0.5, layer_widths=[6, 18,
489     ↪ 12], dropout_rate=0.5, num_classes=10)
490 criterion1 = nn.CrossEntropyLoss()
491 optimizer1 = optim.SGD(model1.parameters(),
492     ↪ lr=0.01, momentum=0.0)
493
494 model2 = DenseNet(growth_rate=24,
495     ↪ compression_factor=0.5, layer_widths=[6, 18,
496     ↪ 12], num_classes=10)
497 criterion2 = nn.CrossEntropyLoss()
498 optimizer2 = optim.SGD(model2.parameters(),
499     ↪ lr=0.01, momentum=0.0)
500
501 model3 = DenseNet(growth_rate=24,
502     ↪ compression_factor=0.5, layer_widths=[6, 18,
503     ↪ 12], dropout_rate=0.5, num_classes=10)
504 criterion3 = nn.CrossEntropyLoss()
505 optimizer3 = optim.SGD(model3.parameters(), lr=0.1,
506     ↪ momentum=0.0)
507
508 model4 = DenseNet(growth_rate=24,
509     ↪ compression_factor=0.5, layer_widths=[6, 18,
510     ↪ 12], num_classes=10)
511 criterion4 = nn.CrossEntropyLoss()
512 optimizer4 = optim.SGD(model4.parameters(), lr=0.1,
513     ↪ momentum=0.0)
514
515 model5 = DenseNet(growth_rate=24,
516     ↪ compression_factor=0.5, layer_widths=[6, 18,
517     ↪ 12], num_classes=10)
518 criterion5 = nn.CrossEntropyLoss()
519 optimizer5 = optim.SGD(model5.parameters(), lr=0.1,
520     ↪ momentum=0.9)
521
522 model6 = DenseNet(growth_rate=24,
523     ↪ compression_factor=0.5, layer_widths=[6, 18,
524     ↪ 12], num_classes=10, version='BC')
525 criterion6 = nn.CrossEntropyLoss()
526 optimizer6 = optim.SGD(model6.parameters(), lr=0.1,
527     ↪ momentum=0.0)
528
529 model7 = DenseNet(growth_rate=24,
530     ↪ compression_factor=0.5, layer_widths=[12, 24,
531     ↪ 18], num_classes=10)
532 criterion7 = nn.CrossEntropyLoss()
533 optimizer7 = optim.SGD(model7.parameters(), lr=0.1,
534     ↪ momentum=0)
535
536 transform = transforms.Compose([
537     transforms.ToTensor(),
538     transforms.Normalize((0.5, 0.5, 0.5), (0.5,
539     ↪ 0.5, 0.5))
540 ])
541
542 trainset = datasets.CIFAR10(root='./data',
543     ↪ train=True, download=True, transform=transform)
544
545 # Set the seed for reproducibility
546 torch.manual_seed(0)
547
548 # Size of trainset
549 trainset_size = len(trainset)
550
551 # Calculate validation set size (10% of trainset
552     ↪ size)
553 valset_size = int(0.1 * trainset_size)
554
555 # Calculate train set size by subtracting
556     ↪ validation set size from total trainset size
557 trainset_size = trainset_size - valset_size
558
559 # Splitting the dataset
560 train_subset, val_subset = random_split(trainset,
561     ↪ [trainset_size, valset_size])
562
563 # Create data loaders for each set
564 trainloader = DataLoader(train_subset,
565     ↪ batch_size=64, shuffle=True)
566 valloader = DataLoader(val_subset, batch_size=64,
567     ↪ shuffle=False)
568
569 testset = datasets.CIFAR10(root='./data',
570     ↪ train=False, download=True, transform=transform)
571 testloader = DataLoader(testset, batch_size=64,
572     ↪ shuffle=False)
573
574 # Defining model, criterions, and optimizers.
575 models = [model1, model2, model3, model4, model5,
576     ↪ model6, model7]
577 criterions = [criterion1, criterion2, criterion3,
578     ↪ criterion4, criterion5, criterion6, criterion7]
579 optimizers = [optimizer1, optimizer2, optimizer3,
580     ↪ optimizer4, optimizer5, optimizer6, optimizer7]
581
582
583 for i in range(len(models)):
584     model = models[i]
585     criterion = criterions[i]
586     optimizer = optimizers[i]
587
588     # Prints the summary of the model. Gives
589     # information about the model parameters,
590     # layers etc.
591     summary(model, input_size=(3, 32, 32))
592
593     # Train the model according to the function
594     # called "train_model"
595     trained_model, train_loss_values,
596     ↪ train_accuracy_values, val_loss_values,
597     ↪ val_accuracy_values = train_model(model,
598     ↪ trainloader, valloader, optimizer,
599     ↪ criterion, device, 60, 7, i+1)
600
601     # Saves the pretrained model if needed later
602     torch.save(trained_model.state_dict(),
603     ↪ f'model_{i+1}_state.pth')
604
605     # Calculates accuracy scores and predicted
606     # labels.
607     train_acc, train_pred = cal_acc(trained_model,
608     ↪ trainloader, device)
609     val_acc, val_pred = cal_acc(trained_model,
610     ↪ valloader, device)
611     test_acc, test_pred = cal_acc(trained_model,
612     ↪ testloader, device)
613
614     print(f'train accuracy: {train_acc}\nvalidation
615     ↪ accuracy: {val_acc}\ntest accuracy:
616     ↪ {test_acc}'[])
617
618     # Calculates confusion matrix for the train,
619     # validation and test set.
620

```

```

573     train_cm =
574         calculate_confusion_matrix(trained_model,
575             trainloader, device)
576         val_cm =
577             calculate_confusion_matrix(trained_model,
578                 valloader, device)
579         test_cm =
580             calculate_confusion_matrix(trained_model,
581                 testloader, device)
582
583     # Plots confusion matrices, loss and accuracy
584     figures.
585     plot_all(train_cm, val_cm, test_cm,
586         train_loss_values, train_accuracy_values,
587         val_accuracy_values, val_loss_values,
588         len(train_accuracy_values), i+1, show=False)
589

```

E. Inception Python Implementation

```

1 import os
2 from argparse import ArgumentParser
3
4 import torch
5
6 from pytorch_lightning import Trainer,
6    seed_everything
7 from pytorch_lightning.callbacks import
7    ModelCheckpoint
8 from pytorch_lightning.loggers import WandbLogger,
8    TensorBoardLogger
9
10 from data import CIFAR10Data
11 from module import CIFAR10Module
12
13
14 def main(args):
15
16     seed_everything(0)
17     os.environ["CUDA_VISIBLE_DEVICES"] = args.gpu_ids
18
19
20     logger = TensorBoardLogger("cifar10",
21         name="inception")
22
23     checkpoint = ModelCheckpoint(monitor="acc/val",
24         mode="max", save_last=False)
25
26     trainer = Trainer(
27         fast_dev_run=bool(args.dev),
28         logger=logger if not bool(args.dev) +
29             args.test_phase) else None,
30         gpus=-1,
31         deterministic=True,
32         weights_summary=None,
33         log_every_n_steps=50,
34         check_val_every_n_epoch=1,
35         max_epochs=args.max_epochs,
36         checkpoint_callback=checkpoint,
37         precision=args.precision,
38     )
39
40     model = CIFAR10Module(args)
41     data = CIFAR10Data(args)
42
43     if bool(args.pretrained):
44         state_dict = os.path.join("state_dicts",
45             "inception" + ".pt")
46         model.model.load_state_dict(torch.load(stat
47             e_dict))
48
49     if bool(args.test_phase):
50

```

```

        trainer.test(model, data.test_dataloader())
51     else:
52         trainer.fit(model, data)
53         trainer.test()
54
55     model.save_metrics()
56
57
58 if __name__ == "__main__":
59     parser = ArgumentParser()
60
61     # PROGRAM level args
62     parser.add_argument("--data_dir", type=str,
63         default="data/cifar10")
64     parser.add_argument("--download_weights",
65         type=int, default=0, choices=[0, 1])
66     parser.add_argument("--test_phase", type=int,
67         default=0, choices=[0, 1])
68     parser.add_argument("--dev", type=int,
69         default=0, choices=[0, 1])
70     parser.add_argument("--logger", type=str,
71         default="tensorboard",
72         choices=["tensorboard", "wandb"])
73
74     # TRAINER args
75     parser.add_argument("--pretrained", type=int,
76         default=0, choices=[0, 1])
77
78     parser.add_argument("--precision", type=int,
79         default=32, choices=[16, 32])
80     parser.add_argument("--batch_size", type=int,
81         default=256)
82     parser.add_argument("--max_epochs", type=int,
83         default=50)
84     parser.add_argument("--num_workers", type=int,
85         default=8)
86     parser.add_argument("--gpu_id", type=str,
87         default="0")
88
89     parser.add_argument("--learning_rate",
90         type=float, default=1e-2)
91     parser.add_argument("--weight_decay",
92         type=float, default=1e-2)
93
94     parser.add_argument("--aux_branches",
95         type=bool, default=False)
96     parser.add_argument("--label_smoothing",
97         type=bool, default=False)
98
99     args = parser.parse_args()
100    main(args)
101

```

```

from typing import Any, List
import pytorch_lightning as pl
import torch
from torchmetrics import Accuracy
import numpy as np
import pickle

from label_smoothing_loss import LabelSmoothingLoss
from model.inception import InceptionModel

from schduler import WarmupCosineLR

```

```

17 class CIFAR10Module(pl.LightningModule):
18     def __init__(self, hparams):
19         super().__init__()
20         self.hparams = hparams
21
22         self.label_smoothing =
23             ↪ self.hparams.label_smoothing
24         self.aux_branches =
25             ↪ self.hparams.aux_branches
26
27         #Loss function selection, with label
28             ↪ smoothing or without.
29         if self.label_smoothing:
30             self.criterion = LabelSmoothingLoss(10)
31         else:
32             self.criterion =
33                 ↪ torch.nn.CrossEntropyLoss()
34
35         self.accuracy = Accuracy()
36         self.accuracy2 = Accuracy(top_k=2)
37         self.accuracy3 = Accuracy(top_k=3)
38
39         self.model = InceptionModel(aux_logits=self.hparams.aux_branches)
40
41         self.epoch_metrics_train = []
42         self.epoch_metrics_val = []
43
44         self.cur_epoch_train_metrics = []
45         self.cur_epoch_val_metrics = []
46
47         self.cur_train_preds = []
48         self.cur_val_preds = []
49         self.cur_test_preds = []
50
51         self.train_preds = []
52         self.val_preds = []
53         self.test_preds = []
54
55
56     def forward(self, batch):
57         images, labels = batch
58         predictions = self.model(images)
59
60         #This is to handle multiple outputs if
61             ↪ auxiliary branches exist.
62         if self.training and self.aux_branches:
63             loss_main =
64                 ↪ self.criterion(predictions[0],
65                     ↪ labels)
66             loss_aux_2 =
67                 ↪ self.criterion(predictions[1],
68                     ↪ labels)
69             loss_aux_1 =
70                 ↪ self.criterion(predictions[2],
71                     ↪ labels)
72
73
74             accuracy =
75                 ↪ self.accuracy(predictions[0],
76                     ↪ labels)
77             accuracy2 =
78                 ↪ self.accuracy2(predictions[0],
79                     ↪ labels)
80             accuracy3 =
81                 ↪ self.accuracy3(predictions[0],
82                     ↪ labels)
83
84             predictions = predictions[0]
85
86         else:
87
88             loss_main = self.criterion(predictions,
89                 ↪ labels)
89             loss_aux_2 = 0
90             loss_aux_1 = 0
91
92             accuracy = self.accuracy(predictions,
93                 ↪ labels)
94             accuracy2 = self.accuracy2(predictions,
95                 ↪ labels)
96             accuracy3 = self.accuracy3(predictions,
97                 ↪ labels)
98
99             predictions = predictions
100
101
102
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125
126
127
128
129
130
131
132
133
134
135
136
137
138
139
140
141

```

loss_main = self.criterion(predictions,
 ↪ labels)
loss_aux_2 = 0
loss_aux_1 = 0

accuracy = self.accuracy(predictions,
 ↪ labels)
accuracy2 = self.accuracy2(predictions,
 ↪ labels)
accuracy3 = self.accuracy3(predictions,
 ↪ labels)

predictions = predictions

#The losses from the auxiliary branches are
 ↪ added with discount weights, the values
 ↪ are 0 if they do not exist.
loss = loss_main + 0.3 * loss_aux_2 + 0.3 *
 ↪ loss_aux_1

return loss, accuracy * 100, accuracy2 *
 ↪ 100, accuracy3 * 100, predictions

```

def on_train_epoch_start(self):
    self.cur_train_preds = []
    self.cur_val_preds = []
    self.cur_test_preds = []

def training_step(self, batch, batch_nb):
    loss, accuracy, acc2, acc3, preds =
        ↪ self.forward(batch)
    self.log("loss/train", loss)
    self.log("acc/train", accuracy)

    metrics = {
        'loss':loss,
        'acc1':accuracy,
        'acc2':acc2,
        'acc3':acc3
    }

    self.cur_epoch_train_metrics.append(metrics)

    pred_probs = preds.cpu().detach().numpy()
    preds = np.argmax(pred_probs, axis=1)

    outs = {
        'preds':preds,
        'pred_probs':pred_probs,
        'labels':batch[1]
    }

    self.cur_train_preds.append(outs)

    return loss

def validation_step(self, batch, batch_nb):
    loss, accuracy, acc2, acc3, preds =
        ↪ self.forward(batch)
    self.log("loss/val", loss)
    self.log("acc/val", accuracy)

    metrics = {
        'loss':loss,
        'acc1':accuracy,
        'acc2':acc2,
        'acc3':acc3
    }

```



```

282     pred_probs = out['pred_probs']           345     self.log("acc/test", accuracy)
283     labels =                                346
284     ↪  out['labels'].cpu().detach().numpy() 347
285     preds_all = np.concatenate((preds_all, 348
286     ↪  preds))                            349
287     pred_probs_all =                      350
288     ↪  np.concatenate((pred_probs_all, 351
289     ↪  pred_probs))                      352
290     labels_all =                          353
291     ↪  np.concatenate((labels_all, labels)) 354
292                                         355
293                                         356
294                                         357
295     out = {                               358
296         'preds':preds_all,                 359
297         'pred_probs':pred_probs_all,       360
298         'labels':labels_all             361
299     }                                     362
300                                         363
301     self.test_preds = out                364
302                                         365
303                                         366
304                                         367
305 def get_metrics(self):                  368
306
307     train_metrics = self.epoch_metrics_train 369
308     val_metrics = self.epoch_metrics_val[1:] 370
309
310     train_outs = self.train_preds          371
311     val_outs = self.val_preds            372
312     test_outs = self.test_preds          373
313
314
315     return train_metrics, val_metrics,    374
316     ↪  train_outs, val_outs, test_outs    375
317
318 def save_metrics(self):                376
319
320     train_metrics, val_metrics, train_outs, 377
321     ↪  val_outs, test_outs = self.get_metrics()
322
323     save = {                           378
324         'train_metrics':train_metrics,   379
325         'val_metrics':val_metrics,     380
326         'train_outs':train_outs,       381
327         'val_outs':val_outs,          382
328         'test_outs':test_outs        383
329     }
330
331     with open('metrics/saved_metrics.pkl', 384
332     ↪  'wb') as f:
333         pickle.dump(save, f)
334
335 def load_metrics():                    385
336
337     with open('metrics/saved_metrics.pkl', 386
338     ↪  'rb') as f:
339         loaded_metrics = pickle.load(f)
340
341
342     return loaded_metrics['train_metrics'], 387
343     ↪  loaded_metrics['val_metrics'],      388
344     ↪  loaded_metrics['train_outs'],       389
345     ↪  loaded_metrics['val_outs'],        390
346     ↪  loaded_metrics['test_outs']
347
348 def test_step(self, batch, batch_nb):  391
349     loss, accuracy, acc2, acc3, preds = 392
350     ↪  self.forward(batch)
351
352
353                                         393
354                                         395
355                                         396
356                                         397
357                                         398
358                                         399
359                                         400
360                                         401
361                                         402
362                                         403
363                                         404
364                                         405
365                                         406
366                                         407
367                                         408
368                                         409
369                                         410
370                                         411
371                                         412
372                                         413
373                                         414
374                                         415
375                                         416
376                                         417
377                                         418
378                                         419
379                                         420
380                                         421
381                                         422
382                                         423
383                                         424
384                                         425
385                                         426
386                                         427
387                                         428
388                                         429
389                                         430
390                                         431
391                                         432
392                                         433
393                                         434
394                                         435
395                                         436
396                                         437
397                                         438
398                                         439
399                                         440
400                                         441
401                                         442
402                                         443
403                                         444
404                                         445
405                                         446
406                                         447
407                                         448
408                                         449
409                                         450
410                                         451
411                                         452
412                                         453
413                                         454
414                                         455
415                                         456
416                                         457
417                                         458
418                                         459
419                                         460
420                                         461
421                                         462
422                                         463
423                                         464
424                                         465
425                                         466
426                                         467
427                                         468
428                                         469
429                                         470
430                                         471
431                                         472
432                                         473
433                                         474
434                                         475
435                                         476
436                                         477
437                                         478
438                                         479
439                                         480
440                                         481
441                                         482
442                                         483
443                                         484
444                                         485
445                                         486
446                                         487
447                                         488
448                                         489
449                                         490
450                                         491
451                                         492
452                                         493
453                                         494
454                                         495
455                                         496
456                                         497
457                                         498
458                                         499
459                                         500
460                                         501
461                                         502
462                                         503
463                                         504
464                                         505
465                                         506
466                                         507
467                                         508
468                                         509
469                                         510
470                                         511
471                                         512
472                                         513
473                                         514
474                                         515
475                                         516
476                                         517
477                                         518
478                                         519
479                                         520
480                                         521
481                                         522
482                                         523
483                                         524
484                                         525
485                                         526
486                                         527
487                                         528
488                                         529
489                                         530
490                                         531
491                                         532
492                                         533
493                                         534
494                                         535
495                                         536
496                                         537
497                                         538
498                                         539
499                                         540
500                                         541
501                                         542
502                                         543
503                                         544
504                                         545
505                                         546
506                                         547
507                                         548
508                                         549
509                                         550
510                                         551
511                                         552
512                                         553
513                                         554
514                                         555
515                                         556
516                                         557
517                                         558
518                                         559
519                                         560
520                                         561
521                                         562
522                                         563
523                                         564
524                                         565
525                                         566
526                                         567
527                                         568
528                                         569
529                                         570
530                                         571
531                                         572
532                                         573
533                                         574
534                                         575
535                                         576
536                                         577
537                                         578
538                                         579
539                                         580
540                                         581
541                                         582
542                                         583
543                                         584
544                                         585
545                                         586
546                                         587
547                                         588
548                                         589
549                                         590
550                                         591
551                                         592
552                                         593
553                                         594
554                                         595
555                                         596
556                                         597
557                                         598
558                                         599
559                                         600
560                                         601
561                                         602
562                                         603
563                                         604
564                                         605
565                                         606
566                                         607
567                                         608
568                                         609
569                                         610
570                                         611
571                                         612
572                                         613
573                                         614
574                                         615
575                                         616
576                                         617
577                                         618
578                                         619
579                                         620
580                                         621
581                                         622
582                                         623
583                                         624
584                                         625
585                                         626
586                                         627
587                                         628
588                                         629
589                                         630
590                                         631
591                                         632
592                                         633
593                                         634
594                                         635
595                                         636
596                                         637
597                                         638
598                                         639
599                                         640
600                                         641
601                                         642
602                                         643
603                                         644
604                                         645
605                                         646
606                                         647
607                                         648
608                                         649
609                                         650
610                                         651
611                                         652
612                                         653
613                                         654
614                                         655
615                                         656
616                                         657
617                                         658
618                                         659
619                                         660
620                                         661
621                                         662
622                                         663
623                                         664
624                                         665
625                                         666
626                                         667
627                                         668
628                                         669
629                                         670
630                                         671
631                                         672
632                                         673
633                                         674
634                                         675
635                                         676
636                                         677
637                                         678
638                                         679
639                                         680
640                                         681
641                                         682
642                                         683
643                                         684
644                                         685
645                                         686
646                                         687
647                                         688
648                                         689
649                                         690
650                                         691
651                                         692
652                                         693
653                                         694
654                                         695
655                                         696
656                                         697
657                                         698
658                                         699
659                                         700
660                                         701
661                                         702
662                                         703
663                                         704
664                                         705
665                                         706
666                                         707
667                                         708
668                                         709
669                                         710
670                                         711
671                                         712
672                                         713
673                                         714
674                                         715
675                                         716
676                                         717
677                                         718
678                                         719
679                                         720
680                                         721
681                                         722
682                                         723
683                                         724
684                                         725
685                                         726
686                                         727
687                                         728
688                                         729
689                                         730
690                                         731
691                                         732
692                                         733
693                                         734
694                                         735
695                                         736
696                                         737
697                                         738
698                                         739
699                                         740
700                                         741
701                                         742
702                                         743
703                                         744
704                                         745
705                                         746
706                                         747
707                                         748
708                                         749
709                                         750
710                                         751
711                                         752
712                                         753
713                                         754
714                                         755
715                                         756
716                                         757
717                                         758
718                                         759
719                                         760
720                                         761
721                                         762
722                                         763
723                                         764
724                                         765
725                                         766
726                                         767
727                                         768
728                                         769
729                                         770
730                                         771
731                                         772
732                                         773
733                                         774
734                                         775
735                                         776
736                                         777
737                                         778
738                                         779
739                                         780
740                                         781
741                                         782
742                                         783
743                                         784
744                                         785
745                                         786
746                                         787
747                                         788
748                                         789
749                                         790
750                                         791
751                                         792
752                                         793
753                                         794
754                                         795
755                                         796
756                                         797
757                                         798
758                                         799
759                                         800
760                                         801
761                                         802
762                                         803
763                                         804
764                                         805
765                                         806
766                                         807
767                                         808
768                                         809
769                                         810
770                                         811
771                                         812
772                                         813
773                                         814
774                                         815
775                                         816
776                                         817
777                                         818
778                                         819
779                                         820
780                                         821
781                                         822
782                                         823
783                                         824
784                                         825
785                                         826
786                                         827
787                                         828
788                                         829
789                                         830
790                                         831
791                                         832
792                                         833
793                                         834
794                                         835
795                                         836
796                                         837
797                                         838
798                                         839
799                                         840
800                                         841
801                                         842
802                                         843
803                                         844
804                                         845
805                                         846
806                                         847
807                                         848
808                                         849
809                                         850
810                                         851
811                                         852
812                                         853
813                                         854
814                                         855
815                                         856
816                                         857
817                                         858
818                                         859
819                                         860
820                                         861
821                                         862
822                                         863
823                                         864
824                                         865
825                                         866
826                                         867
827                                         868
828                                         869
829                                         870
830                                         871
831                                         872
832                                         873
833                                         874
834                                         875
835                                         876
836                                         877
837                                         878
838                                         879
839                                         880
840                                         881
841                                         882
842                                         883
843                                         884
844                                         885
845                                         886
846                                         887
847                                         888
848                                         889
849                                         890
850                                         891
851                                         892
852                                         893
853                                         894
854                                         895
855                                         896
856                                         897
857                                         898
858                                         899
859                                         900
860                                         901
861                                         902
862                                         903
863                                         904
864                                         905
865                                         906
866                                         907
867                                         908
868                                         909
869                                         910
870                                         911
871                                         912
872                                         913
873                                         914
874                                         915
875                                         916
876                                         917
877                                         918
878                                         919
879                                         920
880                                         921
881                                         922
882                                         923
883                                         924
884                                         925
885                                         926
886                                         927
887                                         928
888                                         929
889                                         930
890                                         931
891                                         932
892                                         933
893                                         934
894                                         935
895                                         936
896                                         937
897                                         938
898                                         939
899                                         940
900                                         941
901                                         942
902                                         943
903                                         944
904                                         945
905                                         946
906                                         947
907                                         948
908                                         949
909                                         950
910                                         951
911                                         952
912                                         953
913                                         954
914                                         955
915                                         956
916                                         957
917                                         958
918                                         959
919                                         960
920                                         961
921                                         962
922                                         963
923                                         964
924                                         965
925                                         966
926                                         967
927                                         968
928                                         969
929                                         970
930                                         971
931                                         972
932                                         973
933                                         974
934                                         975
935                                         976
936                                         977
937                                         978
938                                         979
939                                         980
940                                         981
941                                         982
942                                         983
943                                         984
944                                         985
945                                         986
946                                         987
947                                         988
948                                         989
949                                         990
950                                         991
951                                         992
952                                         993
953                                         994
954                                         995
955                                         996
956                                         997
957                                         998
958                                         999
959                                         1000

```

```

32     loss_train = train_metrics[i]['loss'].cpu().detach()
33         ↪ .numpy().item()
34     loss_val = val_metrics[i]['loss'].cpu().detach()
35         ↪ .numpy().item()
36
37     index.append(i)
38     train_losses.append(loss_train)
39     val_losses.append(loss_val)
40
41     axs[0].plot(index, train_losses)
42     axs[0].plot(index, val_losses)
43     axs[0].set_xlabel('Epoch', ylabel='Average Loss')
44     axs[0].set_title('Average Loss vs Epoch')
45     axs[0].legend(['Train', 'Validation'])
46
47 ######
48
49 # Train Accuracy per Epoch
50
51 train_accuracy1 = []
52 train_accuracy2 = []
53 train_accuracy3 = []
54 index = []
55 for i in range(len(train_metrics)):
56     acc1 = train_metrics[i]['acc1'].cpu().detach()
57         ↪ .numpy().item()
58     acc2 = train_metrics[i]['acc2'].cpu().detach()
59         ↪ .numpy().item()
60     acc3 = train_metrics[i]['acc3'].cpu().detach()
61         ↪ .numpy().item()
62
63     index.append(i)
64     train_accuracy1.append(acc1)
65     train_accuracy2.append(acc2)
66     train_accuracy3.append(acc3)
67
68 axs[1].plot(index, train_accuracy1)
69 axs[1].plot(index, train_accuracy2)
70 axs[1].plot(index, train_accuracy3)
71 axs[1].set_xlabel('Epoch', ylabel='Accuracy')
72     ↪ .set_title('Train Accuracy vs Epoch')
73     ↪ .legend(['Top-1 Accuracy', 'Top-2 Accuracy',
74             ↪ 'Top-3 Accuracy'])
75
76 ######
77 # Val Accuracy per Epoch
78
79 val_accuracy1 = []
80 val_accuracy2 = []
81 val_accuracy3 = []
82 index = []
83 for i in range(len(val_metrics)):
84     acc1 = val_metrics[i]['acc1'].cpu().detach()
85         ↪ .numpy().item()
86     acc2 = val_metrics[i]['acc2'].cpu().detach()
87         ↪ .numpy().item()
88     acc3 = val_metrics[i]['acc3'].cpu().detach()
89         ↪ .numpy().item()
90
91     index.append(i)
92     val_accuracy1.append(acc1)
93     val_accuracy2.append(acc2)
94     val_accuracy3.append(acc3)
95
96 axs[2].plot(index, val_accuracy1)
97 axs[2].plot(index, val_accuracy2)
98 axs[2].plot(index, val_accuracy3)
99 axs[2].set_xlabel('Epoch', ylabel='Accuracy')
100    ↪ .set_title('Validation Accuracy vs Epoch')

######
101
102
103 #####
104 fig, axs = plt.subplots(1, 3)
105
106 # Train Confusion Matrix
107 cm = confusion_matrix(train_outs['labels'],
108     ↪ train_outs['preds'])
109
110 disp = ConfusionMatrixDisplay(cm)
111 disp.plot(ax=axs[0])
112 disp.ax_.set_title('Train Confusion Matrix')
113
114 #####
115
116 # Val Confusion Matrix
117 cm = confusion_matrix(val_outs['labels'],
118     ↪ val_outs['preds'])
119
120 disp = ConfusionMatrixDisplay(cm)
121 disp.plot(ax=axs[1])
122 disp.ax_.set_title('Validation Confusion Matrix')
123
124 #####
125
126 # Test Confusion Matrix
127 cm = confusion_matrix(test_outs['labels'],
128     ↪ test_outs['preds'])
129
130 disp = ConfusionMatrixDisplay(cm)
131 disp.plot(ax=axs[2])
132 disp.ax_.set_title('Test Confusion Matrix')
133 fig.suptitle("Confusion Metrics for the model")
134 plt.show()

#####
135
136 # Train Accuracy
137 measureAcc1 = Accuracy(top_k=1)
138 measureAcc2 = Accuracy(top_k=2)
139 measureAcc3 = Accuracy(top_k=3)
140
141 train_acc1 = measureAcc1(torch.from_numpy(train_out
142     ↪ s['pred_probs']),
143     ↪ torch.from_numpy(train_outs['labels']).int())
144 train_acc2 = measureAcc2(torch.from_numpy(train_out
145     ↪ s['pred_probs']),
146     ↪ torch.from_numpy(train_outs['labels']).int())
147 train_acc3 = measureAcc3(torch.from_numpy(train_out
148     ↪ s['pred_probs']),
149     ↪ torch.from_numpy(train_outs['labels']).int())
150
151
152 print("[Train set] Top-1 Accuracy: {acc1:.4f},
153     ↪ Top-2 Accuracy: {acc2:.4f}, Top-3 Accuracy:
154     ↪ {acc3:.4f}.".format(acc1=train_acc1*100,
155     ↪ acc2=train_acc2*100, acc3=train_acc3*100))
156
157 #####
158
159
160

```

```

161
162 # Val Accuracy
163
164
165 val_acc1 = measureAcc1(torch.from_numpy(val_outs['p' 32
166   ↪ red_probs']), 33
167   ↪ torch.from_numpy(val_outs['labels']).int()) 34
168 val_acc2 = measureAcc2(torch.from_numpy(val_outs['p' 35
169   ↪ red_probs']), 36
170   ↪ torch.from_numpy(val_outs['labels']).int()) 37
171 val_acc3 = measureAcc3(torch.from_numpy(val_outs['p' 38
172   ↪ red_probs']), 39
173   ↪ torch.from_numpy(val_outs['labels']).int()) 40
174
175 print("[Validation set] Top-1 Accuracy: {acc1:.4f}, 41
176   ↪ Top-2 Accuracy: {acc2:.4f}, Top-3 Accuracy: 42
177   ↪ {acc3:.4f}, ".format(acc1=val_acc1*100, 43
178   ↪ acc2=val_acc2*100, acc3=val_acc3*100)) 44
179 #####
180
181 # Test Accuracy
182
183
184 test_acc1 = measureAcc1(torch.from_numpy(test_outs['p' 45
185   ↪ pred_probs']), 46
186   ↪ torch.from_numpy(test_outs['labels']).int())
187 test_acc2 = measureAcc2(torch.from_numpy(test_outs['p' 47
188   ↪ pred_probs']), 48
189   ↪ torch.from_numpy(test_outs['labels']).int())
190 test_acc3 = measureAcc3(torch.from_numpy(test_outs['p' 49
191   ↪ pred_probs']), 50
192   ↪ torch.from_numpy(test_outs['labels']).int())
193
194 print("[Test set] Top-1 Accuracy: {acc1:.4f}, Top-2 51
195   ↪ Accuracy: {acc2:.4f}, Top-3 Accuracy: 52
196   ↪ {acc3:.4f}, ".format(acc1=test_acc1*100, 53
197   ↪ acc2=test_acc2*100, acc3=test_acc3*100)) 54
198 #####
199
200
201 import math
202 import warnings
203 from typing import List
204
205 from torch.optim import Optimizer
206 from torch.optim.lr_scheduler import _LRScheduler
207
208
209 class WarmupCosineLR(_LRScheduler):
210     def __init__( 63
211         self,
212         optimizer: Optimizer, 64
213         warmup_epochs: int, 65
214         max_epochs: int, 66
215         warmup_start_lr: float = 1e-8, 67
216         eta_min: float = 1e-8, 68
217         last_epoch: int = -1, 69
218     ) -> None:
219
220         self.warmup_epochs = warmup_epochs
221         self.max_epochs = max_epochs
222         self.warmup_start_lr = warmup_start_lr
223         self.eta_min = eta_min
224
225         super(WarmupCosineLR, 70
226             self).__init__(optimizer, last_epoch)
227
228     def get_lr(self) -> List[float]:
229         if not self._get_lr_called_within_step:
230             warnings.warn(
231                 "To get the last learning rate 71
232                 ↪ computed by the scheduler, " 72
233                 "use `get_last_lr()``.", 73
234
235
236         UserWarning,
237         )
238
239     if self.last_epoch == 0:
240         return [self.warmup_start_lr] *
241             len(self.base_lrs)
242     elif self.last_epoch < self.warmup_epochs:
243         return [
244             group["lr"]
245             + (base_lr - self.warmup_start_lr)
246             / (self.warmup_epochs - 1)
247             for base_lr, group in
248             zip(self.base_lrs,
249                 self.optimizer.param_groups)
250         ]
251     elif self.last_epoch == self.warmup_epochs:
252         return self.base_lrs
253     elif (self.last_epoch - 1 -
254         self.max_epochs) % (
255             2 * (self.max_epochs -
256                 self.warmup_epochs)
257         ) == 0:
258         return [
259             group["lr"]
260             + (base_lr - self.eta_min)
261             * (1 - math.cos(math.pi /
262                 (self.max_epochs -
263                     self.warmup_epochs))) /
264             2
265             for base_lr, group in
266             zip(self.base_lrs,
267                 self.optimizer.param_groups)
268         ]
269
270     return [
271         (
272             1
273             + math.cos(
274                 math.pi
275                 * (self.last_epoch -
276                     self.warmup_epochs) /
277                     (self.max_epochs -
278                         self.warmup_epochs))
279         )
280         /
281         (
282             1
283             + math.cos(
284                 math.pi
285                 * (self.last_epoch -
286                     self.warmup_epochs - 1) /
287                     (self.max_epochs -
288                         self.warmup_epochs)
289             )
290         )
291         *
292         (group["lr"] - self.eta_min)
293         + self.eta_min
294         for group in self.optimizer.param_groups
295     ]
296
297     def _get_closed_form_lr(self) -> List[float]:
298         if self.last_epoch < self.warmup_epochs:
299             return [
300                 self.warmup_start_lr
301                 + self.last_epoch
302                 * (base_lr - self.warmup_start_lr)
303                 / (self.warmup_epochs - 1)
304                 for base_lr in self.base_lrs
305             ]
306
307         return [
308             self.eta_min
309             + 0.5
310             * (base_lr - self.eta_min)
311             *
312                 (1
313                 + math.cos(
314

```

```

95         math.pi
96         * (self.last_epoch -
97             ↪ self.warmup_epochs)
98         / (self.max_epochs -
99             ↪ self.warmup_epochs)
100    )
101   ]
```

```

1 import os
2 from collections import namedtuple
3
4 import torch
5 import torch.nn as nn
6 import torch.nn.functional as F
7
8 __all__ = ["InceptionModel", "InceptionModel"]
9
10 _InceptionModelOuputs = namedtuple(
11     "InceptionModelOuputs", ["logits",
12                             ↪ "aux_logits2", "aux_logits1"])
13 )
14
15
16 def InceptionModel(pretrained=False, progress=True,
17                     ↪ device="cpu", aux_logits=False, **kwargs):
18
19     model = InceptionModel(aux_logits=aux_logits)
20     if pretrained:
21         script_dir = os.path.dirname(__file__)
22         state_dict = torch.load(
23             script_dir +
24             "/state_dicts/InceptionModel.pt",
25             map_location=device
26         )
27         model.load_state_dict(state_dict)
28     return model
29
30
31 class InceptionModel(nn.Module):
32
33     def __init__(self, num_classes=10,
34                  ↪ aux_logits=False):
35         super(InceptionModel, self).__init__()
36         self.aux_logits = aux_logits
37
38         self.conv1 = BasicConv2d(3, 192,
39                               ↪ kernel_size=3, stride=1, padding=1)
40
41         self.inception1 = Inception(192, 64, 96,
42                                     ↪ 128, 16, 32, 32)
43         self.inception2 = Inception(256, 128, 128,
44                                     ↪ 192, 32, 96, 64)
45
46         self.maxpool1 = nn.MaxPool2d(3, stride=2,
47                                     ↪ padding=1, ceil_mode=False)
48
49         self.inception3 = Inception(480, 192, 96,
50                                     ↪ 208, 16, 48, 64)
51         self.inception4 = Inception(512, 160, 112,
52                                     ↪ 224, 24, 64, 64)
53         self.inception5 = Inception(512, 128, 128,
54                                     ↪ 256, 24, 64, 64)
55         self.inception6 = Inception(512, 112, 144,
56                                     ↪ 288, 32, 64, 64)
57         self.inception7 = Inception(528, 256, 160,
58                                     ↪ 320, 32, 128, 128)
59
60         self.maxpool2 = nn.MaxPool2d(3, stride=2,
61                                     ↪ padding=1, ceil_mode=False)
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111
112
113
114
```

```

115     self.branch2 = nn.Sequential(
116         BasicConv2d(in_channels, ch3x3red,
117             ↪ kernel_size=1),
118         BasicConv2d(ch3x3red, ch3x3,
119             ↪ kernel_size=3, padding=1),
120     )
121
122     self.branch3 = nn.Sequential(
123         BasicConv2d(in_channels, ch5x5red,
124             ↪ kernel_size=1),
125         BasicConv2d(ch5x5red, ch5x5,
126             ↪ kernel_size=3, padding=1),
127     )
128
129     self.branch4 = nn.Sequential(
130         nn.MaxPool2d(kernel_size=3, stride=1,
131             ↪ padding=1, ceil_mode=True),
132         BasicConv2d(in_channels, pool_proj,
133             ↪ kernel_size=1),
134     )
135
136     def forward(self, x):
137         branch1 = self.branch1(x)
138         branch2 = self.branch2(x)
139         branch3 = self.branch3(x)
140         branch4 = self.branch4(x)
141
142         outputs = [branch1, branch2, branch3,
143             ↪ branch4]
144         return torch.cat(outputs, 1)
145
146 class BasicConv2d(nn.Module):
147     def __init__(self, in_channels, out_channels,
148         ↪ **kwargs):
149         super(BasicConv2d, self).__init__()
150         self.conv = nn.Conv2d(in_channels,
151             ↪ out_channels, bias=False, **kwargs)
152         self.bn = nn.BatchNorm2d(out_channels,
153             ↪ eps=0.001)
154
155     def forward(self, x):
156         x = self.conv(x)
157         x = self.bn(x)
158         return F.relu(x, inplace=True)
159
160 #This is the ending sequence of the Auxiliary
161 #→ Branches.
162 #The outputs are treated as normal model outputs to
163 #→ calculate loss in module.py, expect only in
164 #→ training.
165 class InceptionAux(nn.Module):
166     def __init__(self, in_channels, num_classes):
167         super(InceptionAux, self).__init__()
168         self.conv = BasicConv2d(in_channels, 128,
169             ↪ kernel_size=1)
170
171         self.fc1 = nn.Linear(2048, 1024)
172         self.fc2 = nn.Linear(1024, num_classes)
173
174     def forward(self, x):
175         x = F.adaptive_avg_pool2d(x, (4, 4))
176
177         x = self.conv(x)
178
179         x = x.view(x.size(0), -1)
180
181         x = F.relu(self.fc1(x), inplace=True)
182
183         x = F.dropout(x, 0.7,
184             ↪ training=self.training)
185
186         x = self.fc2(x)
187
188         return x

```

```

1 import os
2 import zipfile
3
4 import pytorch_lightning as pl
5 import requests
6 from torch.utils.data import DataLoader, Subset
7 from torchvision import transforms as T
8 from torchvision.datasets import CIFAR10
9 from tqdm import tqdm
10
11
12 class CIFAR10Data(pl.LightningDataModule):
13     def __init__(self, args):
14         super().__init__()
15         self.hparams = args
16         self.mean = (0.4914, 0.4822, 0.4465)
17         self.std = (0.2471, 0.2435, 0.2616)
18
19
20     def train_dataloader(self):
21         transform = T.Compose(
22             [
23                 T.RandomCrop(32, padding=4),
24                 T.RandomHorizontalFlip(),
25                 T.ToTensor(),
26                 T.Normalize(self.mean, self.std),
27             ]
28         )
29         dataset =
30             ↪ CIFAR10(root=self.hparams.data_dir,
31             ↪ train=True, transform=transform)
32         #subset = Subset(dataset, range(256))
33
34         dataloader = DataLoader(
35             dataset,
36             batch_size=self.hparams.batch_size,
37             num_workers=self.hparams.num_workers,
38             shuffle=True,
39             drop_last=True,
40             pin_memory=True,
41         )
42         return dataloader
43
44     def val_dataloader(self):
45         transform = T.Compose(
46             [
47                 T.ToTensor(),
48                 T.Normalize(self.mean, self.std),
49             ]
50         )
51         dataset =
52             ↪ CIFAR10(root=self.hparams.data_dir,
53             ↪ train=False, transform=transform)
54         #subset = Subset(dataset, range(256))
55
56         dataloader = DataLoader(
57             dataset,
58             batch_size=self.hparams.batch_size,
59             num_workers=self.hparams.num_workers,
60             drop_last=True,
61             pin_memory=True,
62         )
63         return dataloader
64
65     def test_dataloader(self):
66         return self.val_dataloader()

```

```

1 import torch
2 import numpy as np
3 import torch.nn as nn
4 import torch.nn.functional as F
5 from torch.autograd import Variable

```

